

# CSC 422/522

# Computer Vision and

# Pattern Recognition

## Image Processing (Basics)

### 2026 Spring



SOUTH DAKOTA  
STATE UNIVERSITY

# Today

- What is an image?
  - Color
  - Histogram
  - Texture
- How to resize an image?
- OpenCV practice in Colab
  - Read, Show, and Write
  - Access and Modify Pixel Values
  - Access Image Properties
  - Color Space Transformation
  - Crop
  - Histogram Equalization
  - Resize



SOUTH DAKOTA  
STATE UNIVERSITY



SOUTH DAKOTA  
STATE UNIVERSITY



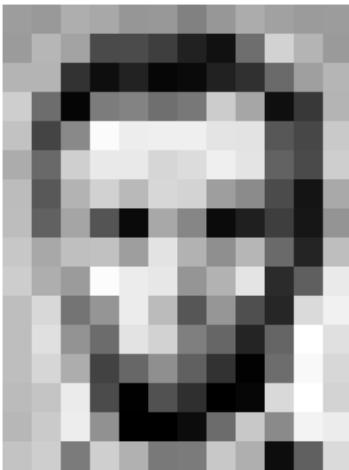
Image taken at Niagara Falls during my graduation trip ☺  
Hope it brings you good luck!

# What is an image?

Credits: some slides of this section are adapted from CSC 492/592 (Spring 2024)

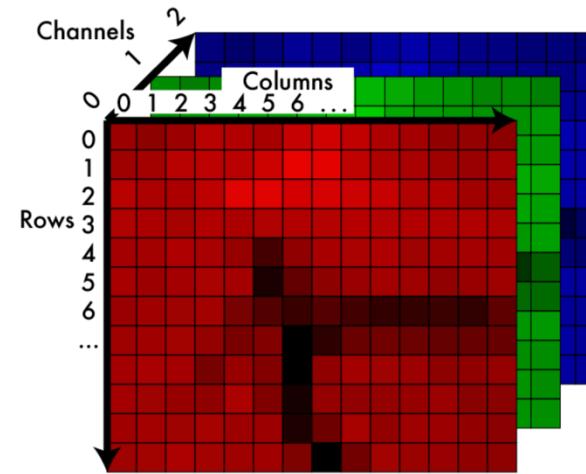
# What is an image?

- We present an image as grids of pixels
- ***Resolution:*** width / height
- ***Channel:*** color / gray



157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	83	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	299	299	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	96	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

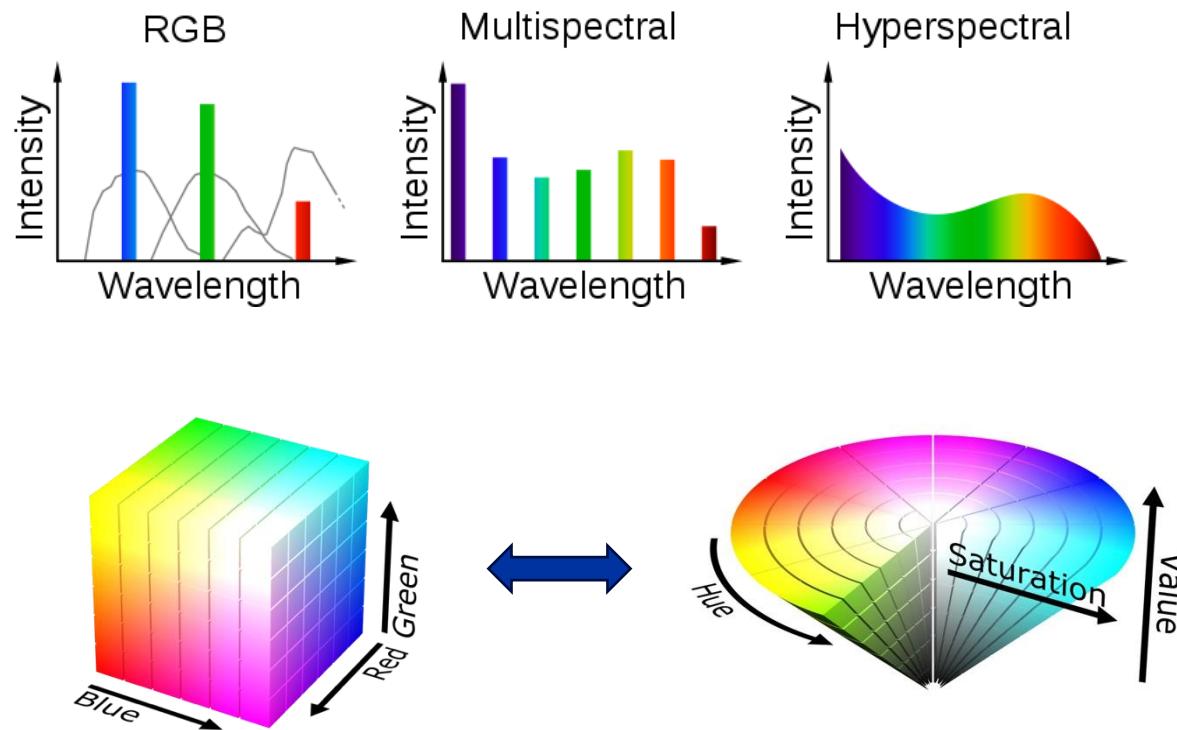
157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	83	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	299	299	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	96	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218



SOUTH DAKOTA  
STATE UNIVERSITY

# What is color?

- Image data can be saved in different color spaces

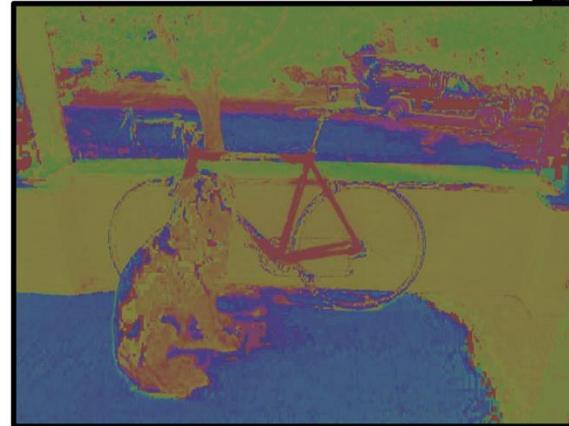


SOUTH DAKOTA  
STATE UNIVERSITY

# Let's Play with Image

- HSV (Hue, Saturation, Value) is generally considered closer to human perception than RGB.

Hue



Saturation



Value

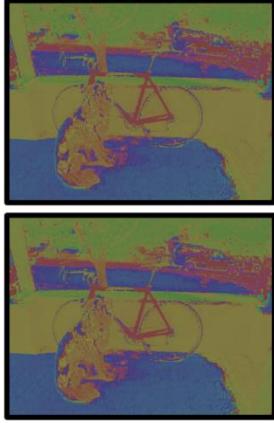


SOUTH DAKOTA  
STATE UNIVERSITY

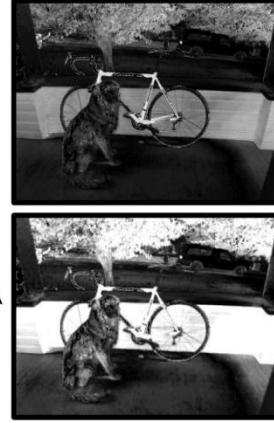
# Let's Play with Image

More saturation = intense colors

H Channel



S Channel



V Channel



2x

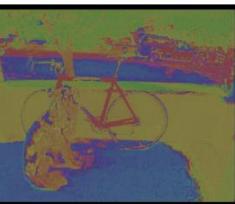
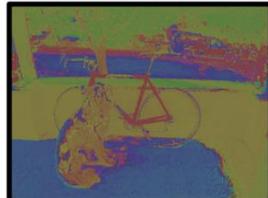


SOUTH DAKOTA  
STATE UNIVERSITY

# Let's Play with Image

More value = lighter image

H Channel



S Channel



V Channel



2x

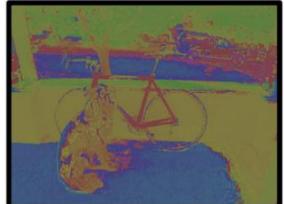


SOUTH DAKOTA  
STATE UNIVERSITY

# Let's Play with Image

Shift hue = shift colors

H Channel



S Channel



V Channel



- .2

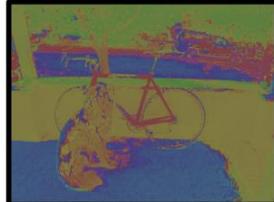


SOUTH DAKOTA  
STATE UNIVERSITY

# Let's Play with Image

Set hue to your favorite color!

H Channel



S Channel



V Channel

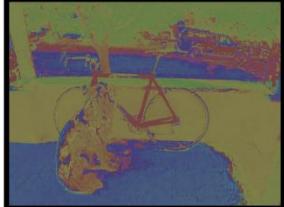


SOUTH DAKOTA  
STATE UNIVERSITY

# Let's Play with Image

Or pattern...

H Channel



S Channel



V Channel

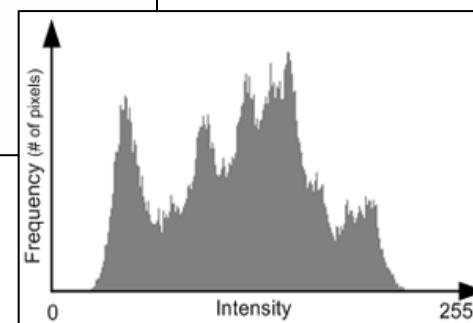


SOUTH DAKOTA  
STATE UNIVERSITY

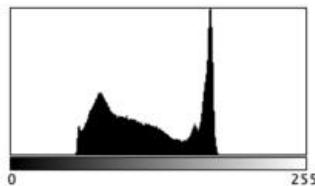
# Color: Histograms

- A histogram of a grayscale image is an array  $H[*]$  of bins, one for each gray tone
- $H[i]$  gives the count of how many pixels of an image have gray tone  $i$
- $P[i]$  (the normalized histogram) gives the percentage of pixels that have gray one  $i$

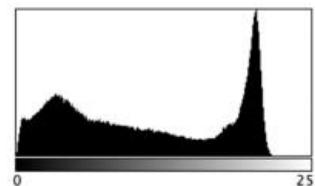
```
// traverse all pixels and accumulate the count of same intensity values
for ( i = 0; i < pixelCount; ++i )
{
    histogram[image[i]]++;
}
```



# Color: Histograms



low contrast



normal contrast



high contrast



SOUTH DAKOTA  
STATE UNIVERSITY

<https://pixelcraft.photo.blog/2022/06/01/the-image-histogram-vs-contrast-and-clipping/#:~:text=Figure%20shows%20an%20example,contrast%20on%20a%20grayscale%20image.&text=The%20histogram%20of%20a%20high,th%20way%20of%20tonal%20contrast.>

# What is a Texture?

- Image texture is a set of metrics to quantify the perceived texture of an image



SOUTH DAKOTA  
STATE UNIVERSITY

# What is a Texture? Local Binary Pattern (LBP)

- For each pixel  $p$ , create an 8-bit number  $b_1, b_2, b_3, b_4, b_5, b_6, b_7, b_8$ , where  $b_i = 0$  if neighbor  $i$  has value less than or equal to  $p$ 's value and 1 otherwise.
- Represent the texture in the image (or a region) by the histogram of these numbers.

10	10	10
0	1	3
40	50	80
50	60	90

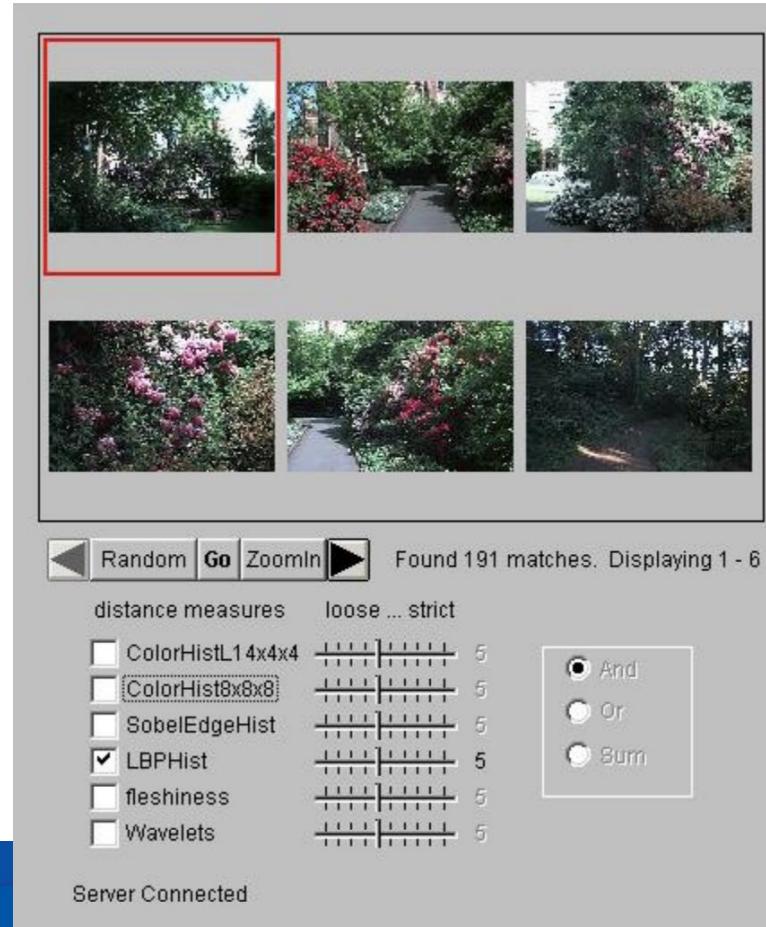


11111100



# What is a Texture? Local Binary Pattern (LBP)

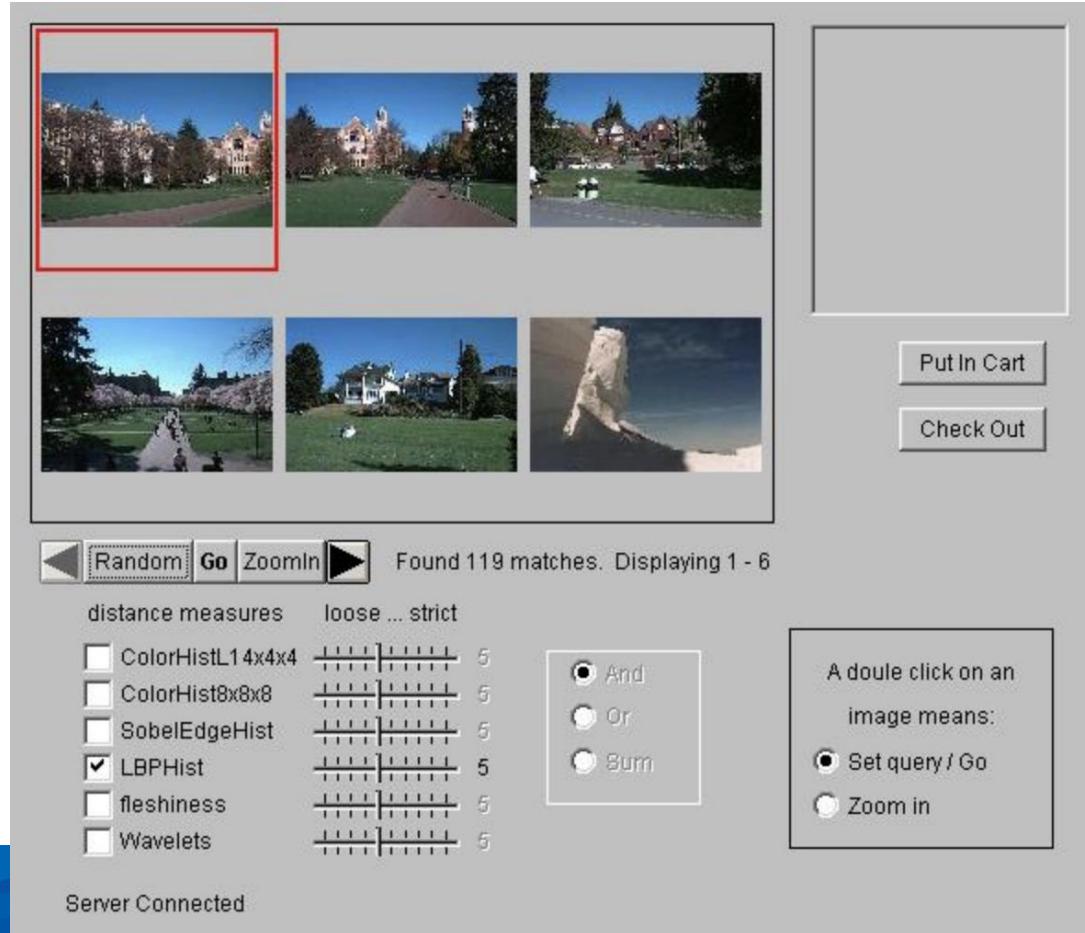
- Example:
  - Image database system retrieves images similar to the query image using LBP texture



SOUTH DAKOTA  
STATE UNIVERSITY

# What is a Texture? Local Binary Pattern (LBP)

- Not always semantically similar images
- But, effective for low-level semantic information
- What if we need more high-level semantics?



SOUTH DAKOTA  
STATE UNIVERSITY



SOUTH DAKOTA  
STATE UNIVERSITY



# How to Resize an Image?

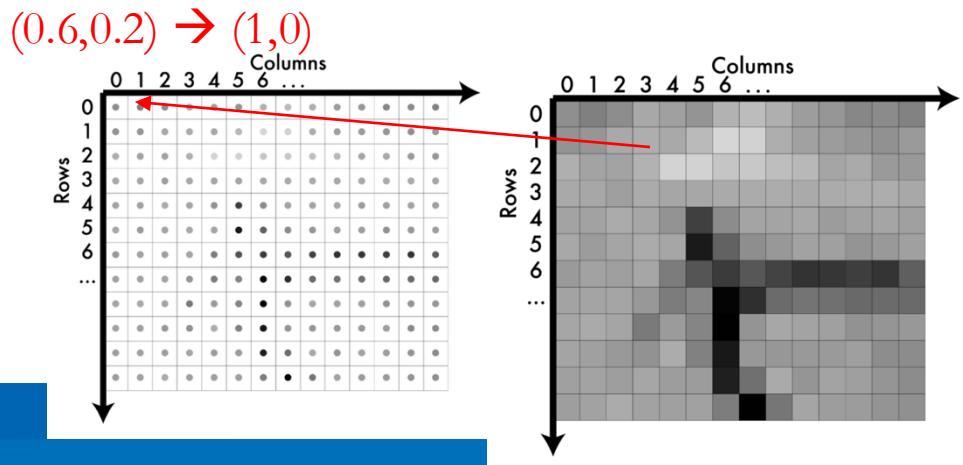
Credits: some slides of this section are adapted from CSC 492/592 (Spring 2024)



Image taken at Niagara Falls during  
my graduation trip ☺

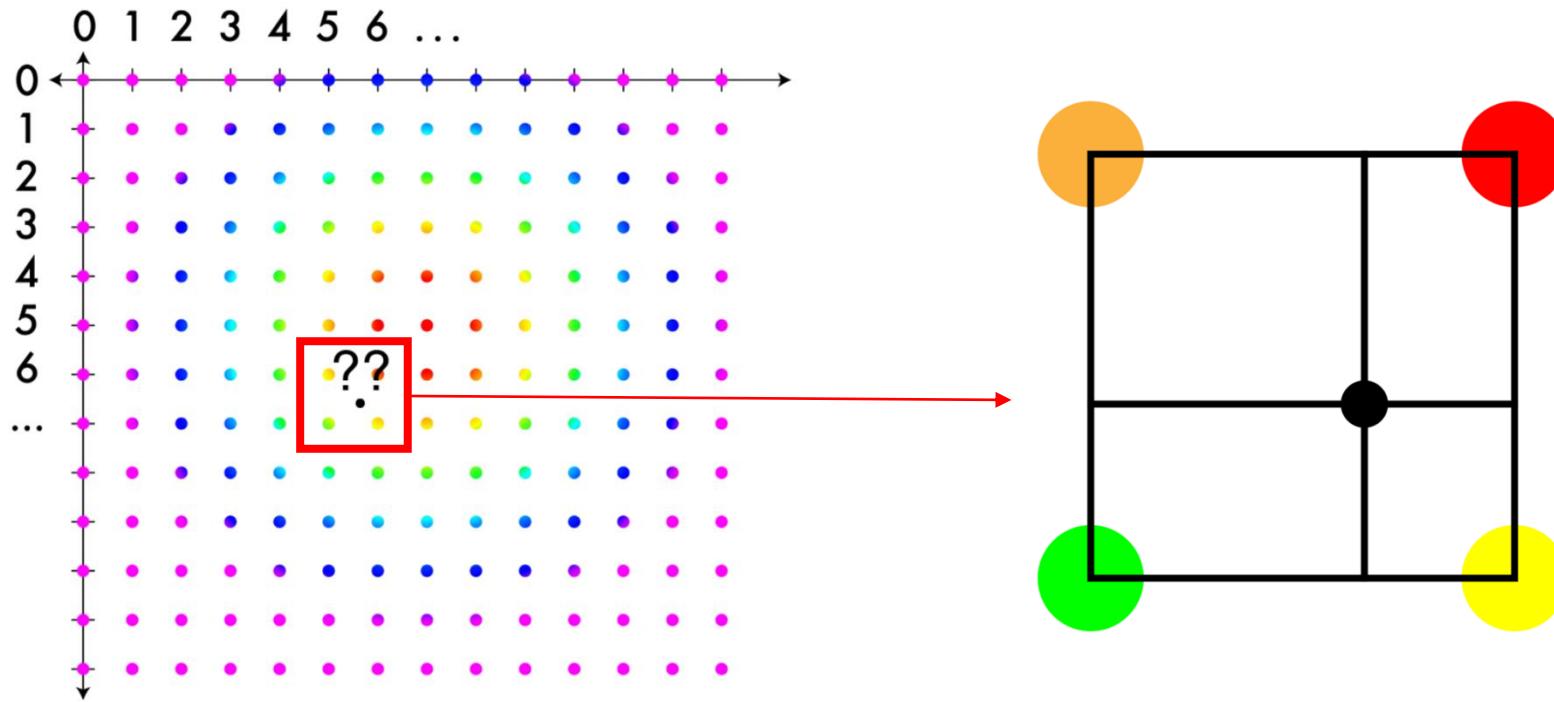
# Method 1: Nearest-Neighbor Interpolation

- $f(x, y) = Im(round\left(\frac{x}{s}\right), round\left(\frac{y}{s}\right))$



SOUTH DAKOTA  
STATE UNIVERSITY

# Method 2: Bilinear Interpolation



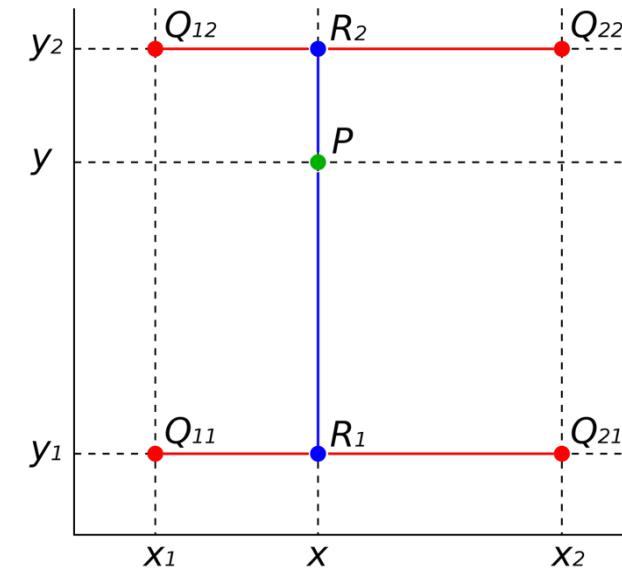
SOUTH DAKOTA  
STATE UNIVERSITY

# Method 2: Bilinear Interpolation

$$R_1: f(x, y_1) = \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21})$$

$$R_2: f(x, y_2) = \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22})$$

$$P: f(x, y) = \frac{y_2 - y}{y_2 - y_1} f(x, y_1) + \frac{y - y_1}{y_2 - y_1} f(x, y_2)$$



SOUTH DAKOTA  
STATE UNIVERSITY

# Method 2: Bilinear Interpolation

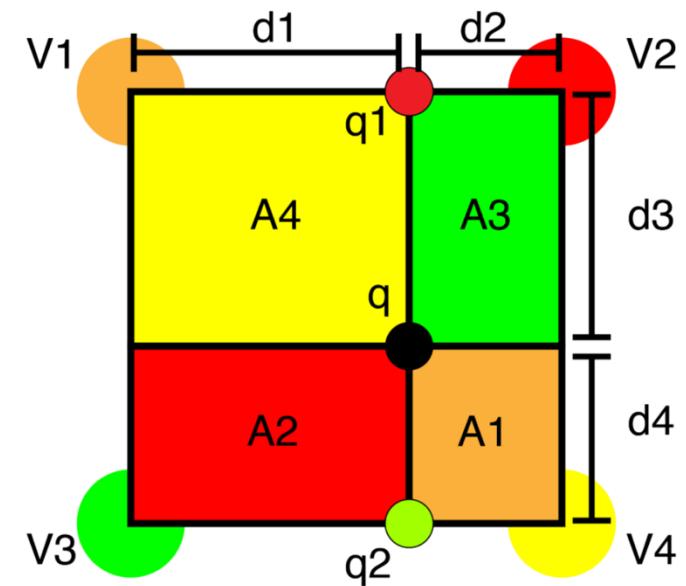
$$q = V1 * A1 + V2 * A2 + V3 * A3 + V4 * A4$$

$$A1 = d2 * d4$$

$$A2 = d1 * d4$$

$$A3 = d2 * d3$$

$$A4 = d1 * d3$$



# Nearest-Neighbor vs. Bilinear

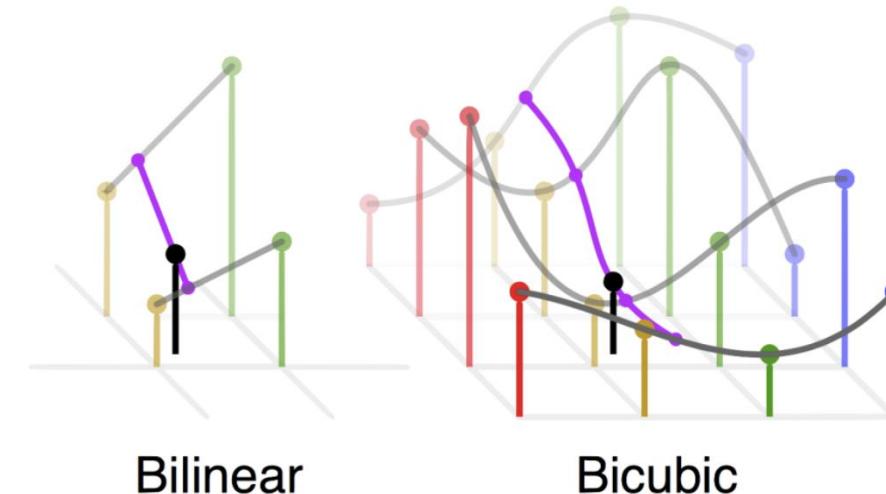
- Smoother than NN
- More complex
- Tradeoff: speed vs. quality



SOUTH DAKOTA  
STATE UNIVERSITY

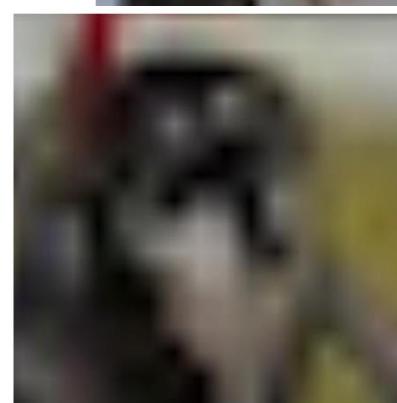
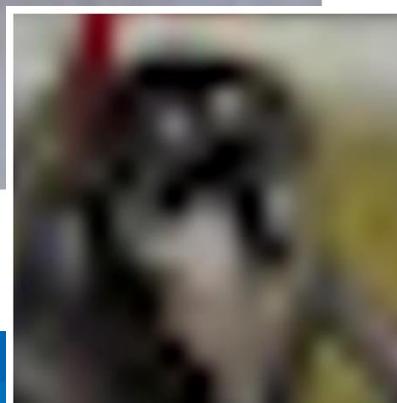
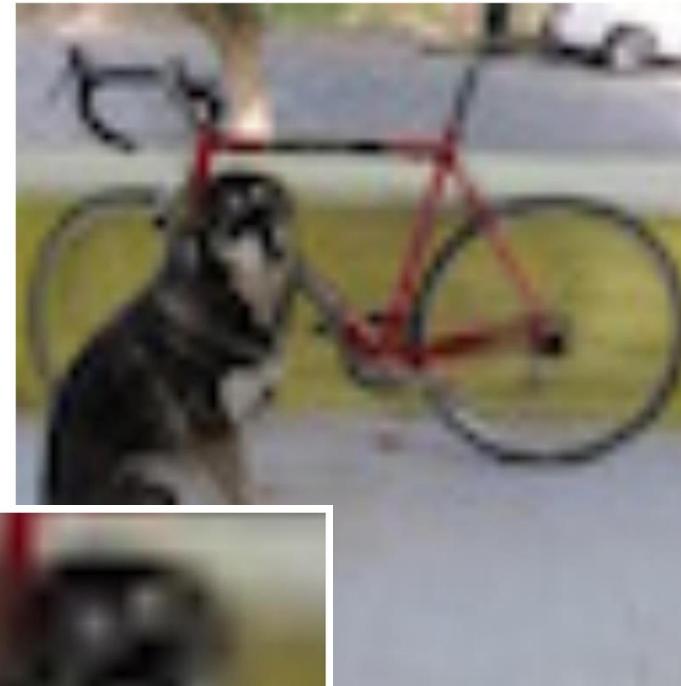
# Method 3: Bicubic Interpolation

- Smoother than bilinear
- 16 nearest neighbors
- Fit 3<sup>rd</sup> order poly:
  - $f(x) = a_0x^0 + a_1x^1 + a_2x^2$
- Interpolate along axis
- Fit another poly to another axis



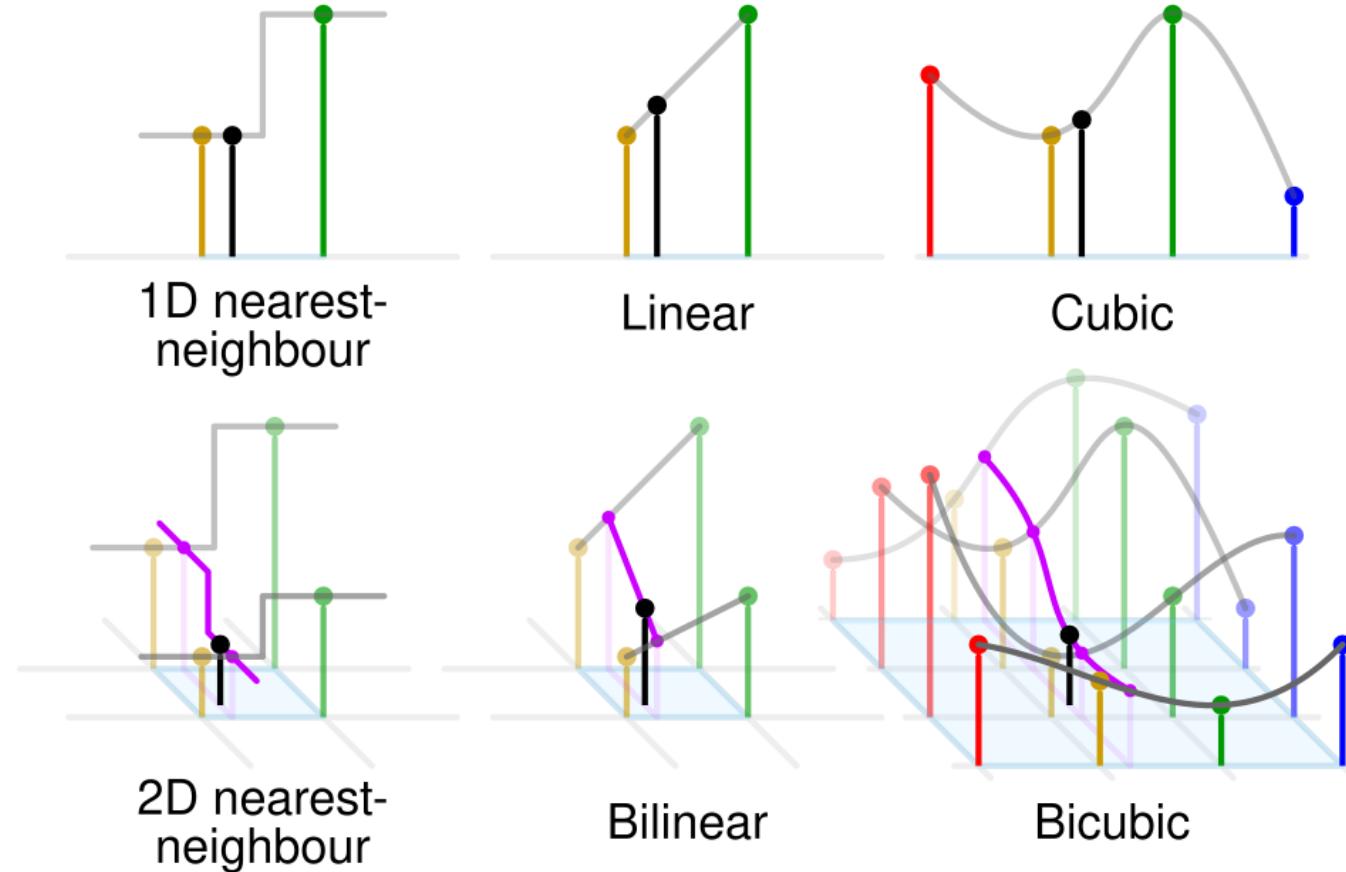
SOUTH DAKOTA  
STATE UNIVERSITY

# Bicubic vs. Bilinear



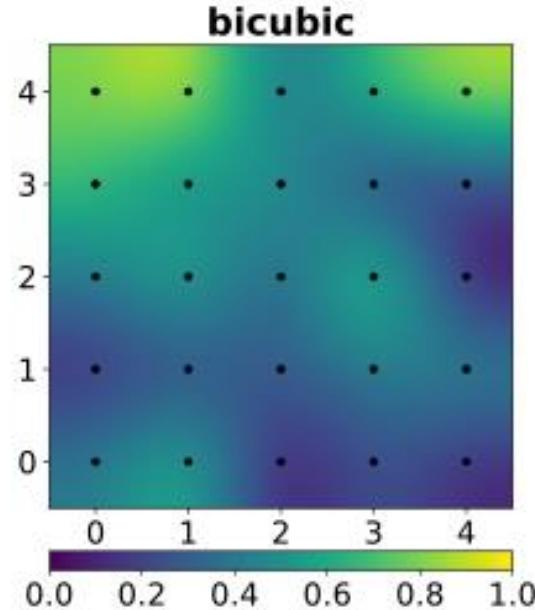
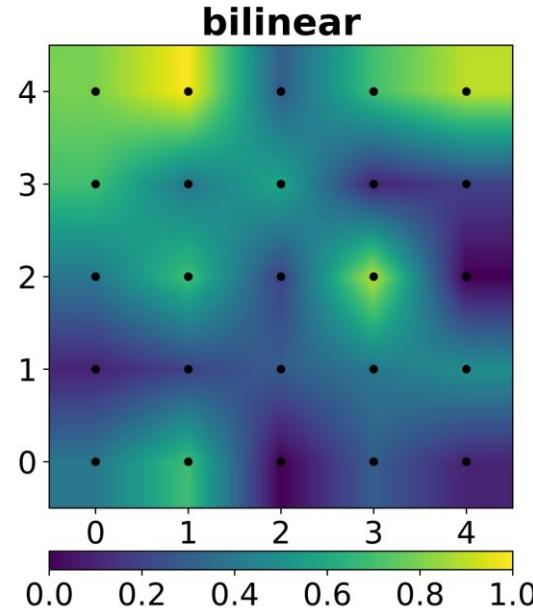
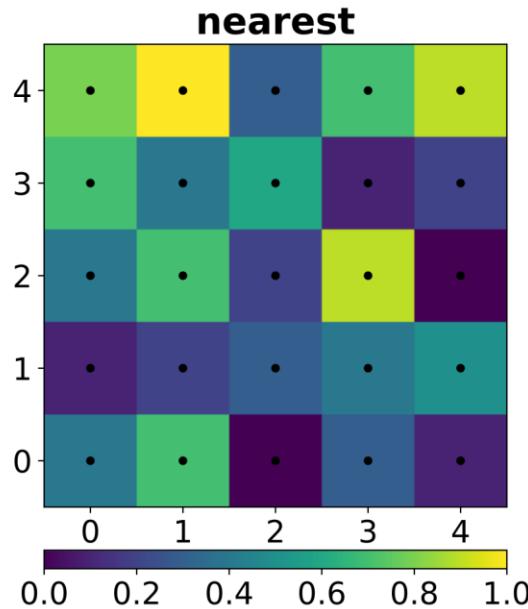
SOUTH DAKOTA  
STATE UNIVERSITY

# Comparison of bicubic interpolation with some 1- and 2-dimensional interpolations



Source:  
[https://en.wikipedia.org  
/wiki/File:Comparison  
of\\_1D\\_and\\_2D\\_interpol  
ation.svg](https://en.wikipedia.org/wiki/File:Comparison_of_1D_and_2D_interpolation.svg)

# Nearest-Neighbor vs. Bilinear vs. Bicubic



Source: [https://en.wikipedia.org/wiki/Bicubic\\_interpolation](https://en.wikipedia.org/wiki/Bicubic_interpolation)



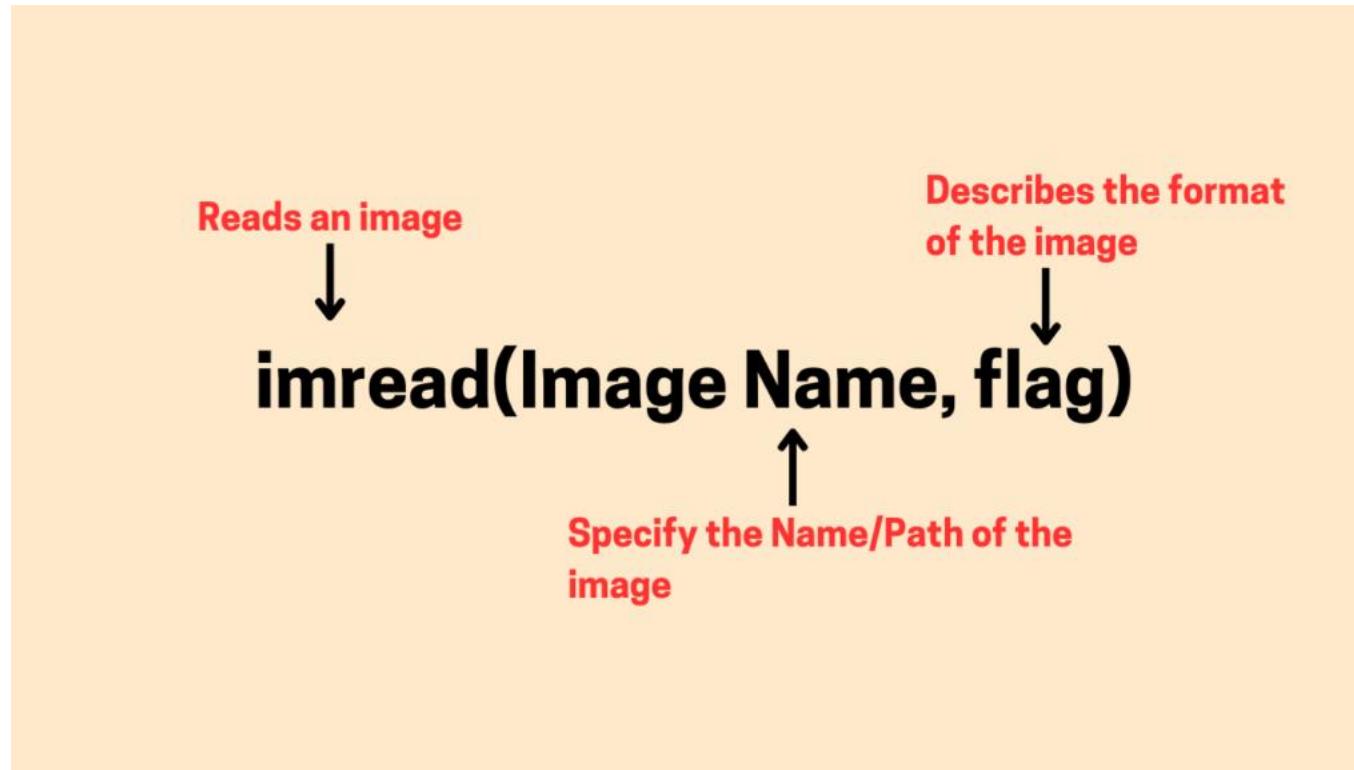
SOUTH DAKOTA  
STATE UNIVERSITY



Image from: <https://opencv.org/blog/read-display-and-write-an-image-using-opencv/>

# OpenCV practice in Colab

# Reading an Image in OpenCV



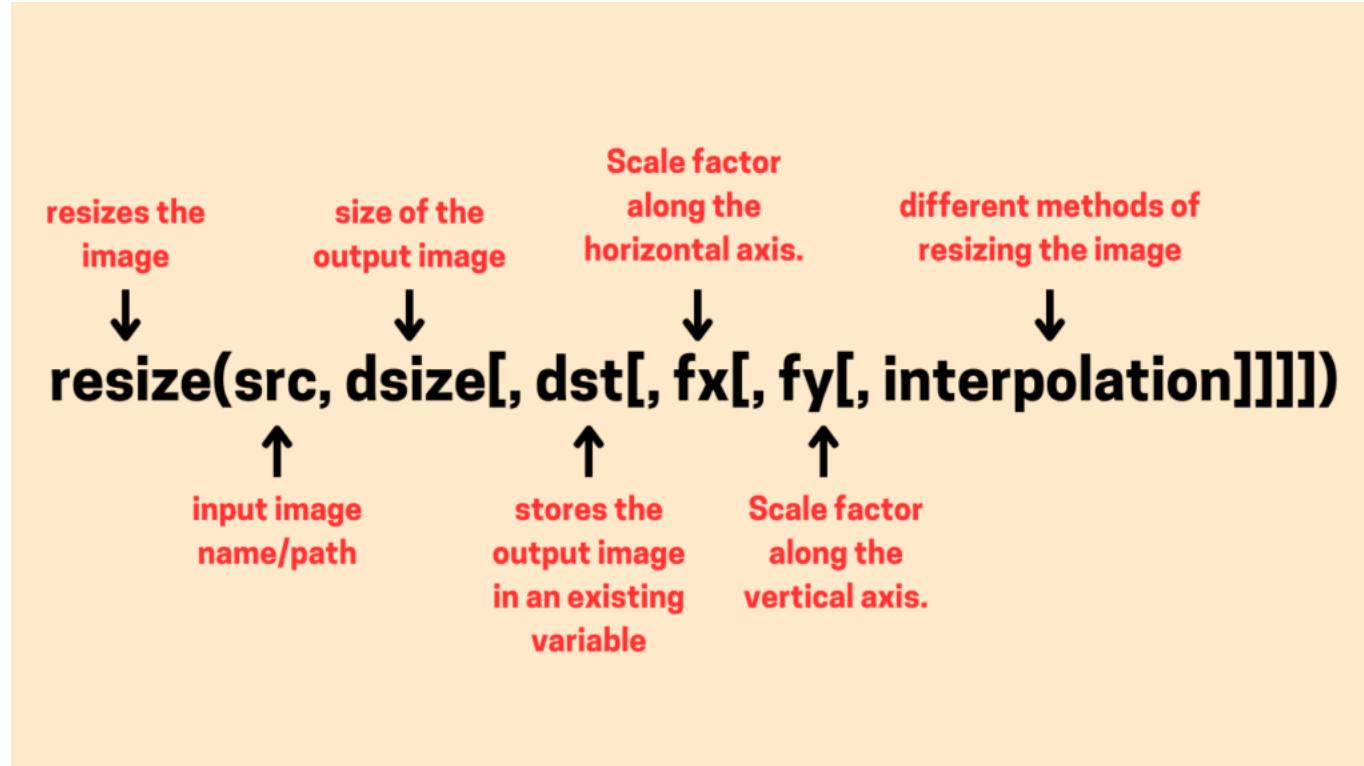
SOUTH DAKOTA  
STATE UNIVERSITY

# Reading an Image in OpenCV

Flags	Alternative Value	Description
<b>IMREAD_UNCHANGED</b>	-1	Loads the image as is, including the alpha channel (if present). Total of 4 channels – (3-BGR & 1-Alpha)
<b>IMREAD_GRAYSCALE</b>	0	Loads the image in grayscale mode (single channel).
<b>IMREAD_COLOR</b>	1	Loads the image in full color (default mode, 3-BGR) ignores alpha channel

<https://opencv.org/blog/read-display-and-write-an-image-using-opencv/>

# Resize an Image in OpenCV



<https://opencv.org/blog/resizing-and-rescaling-images-with-opencv/#h-opencv-function-to-resize-images>

# Resize an Image in OpenCV

Method	Description	Best Used For
INTER_NEAREST	Nearest-neighbor interpolation (fastest, but low quality)	Simple, fast resizing (e.g., pixel art, binary images)
INTER_LINEAR	Bilinear interpolation (default method)	General-purpose resizing (good balance of speed & quality)
INTER_CUBIC	Bicubic interpolation (uses $4 \times 4$ pixel neighborhood)	High-quality upscaling, smoother results
INTER_AREA	Resampling using pixel area relation	Best for <b>shrinking</b> images (avoids aliasing)
INTER_LANCZOS4	Lanczos interpolation using $8 \times 8$ pixel neighborhood	High-quality <b>upscaling &amp; downscaling</b> (preserves fine details)

<https://opencv.org/blog/resizing-and-rescaling-images-with-opencv/#h-opencv-function-to-resize-images>

# Let's Practice More Operations in Colab!

- Link:

[https://colab.research.google.com/drive/1wZt00QnxYgMVjlih16  
RI73IS\\_99UIHQt?usp=sharing](https://colab.research.google.com/drive/1wZt00QnxYgMVjlih16RI73IS_99UIHQt?usp=sharing)



SOUTH DAKOTA  
STATE UNIVERSITY

# Takeaway

- Image properties: color, histogram, texture
- Different interpolation methods to resize an image
- Use OpenCV to perform basic image processing
  - Color space transformation
  - Crop and resize
  - Histogram Equalization



Questions?



SOUTH DAKOTA  
STATE UNIVERSITY