

# CSC 422/522

# Computer Vision and

# Pattern Recognition

Filtering (Convolution) and Feature Detection

2026 Spring



SOUTH DAKOTA  
STATE UNIVERSITY

# Today

- Boundary Effects and Padding
- Binary Image Processing
- Features vs. Pixels
- How to detect edges?



SOUTH DAKOTA  
STATE UNIVERSITY



SOUTH DAKOTA  
STATE UNIVERSITY

# Boundary Effects and Padding

Image source: <https://datahacker.rs/what-is-padding-cnn/>

0	0	0	0	0	0	0	0	0
0	3	3	4	4	7	0	0	0
0	9	7	6	5	8	2	0	0
0	6	5	5	6	9	2	0	0
0	7	1	3	2	7	8	0	0
0	0	3	7	1	8	3	0	0
0	4	0	4	3	2	2	0	0
0	0	0	0	0	0	0	0	0

$6 \times 6 \rightarrow 8 \times 8$

\*

1	0	-1
1	0	-1
1	0	-1

$3 \times 3$

=

-10	-13	1						
-9	3	0						

$6 \times 6$

# Recap: Convolution

$$\begin{array}{|c|c|c|c|c|} \hline 7 & 2 & 3 & 3 & 8 \\ \hline 4 & 5 & 3 & 8 & 4 \\ \hline 3 & 3 & 2 & 8 & 4 \\ \hline 2 & 8 & 7 & 2 & 7 \\ \hline 5 & 4 & 4 & 5 & 4 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 6 & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$$7x1 + 4x1 + 3x1 +  
2x0 + 5x0 + 3x0 +  
3x-1 + 3x-1 + 2x-1  
= 6$$

Animation source: <https://medium.datadriveninvestor.com/convolutional-neural-networks-3b241a5da51e>

# Boundary Effects

$$\begin{array}{|c|c|c|c|c|c|c|c|c|} \hline 45 & 60 & 98 & 127 & 132 & 133 & 137 & 133 \\ \hline 46 & 65 & 98 & 123 & 126 & 128 & 131 & 133 \\ \hline 47 & 65 & 96 & 115 & 119 & 123 & 135 & 137 \\ \hline 47 & 63 & 91 & 107 & 113 & 122 & 138 & 134 \\ \hline 50 & 59 & 80 & 97 & 110 & 123 & 133 & 134 \\ \hline 49 & 53 & 68 & 83 & 97 & 113 & 128 & 133 \\ \hline 50 & 50 & 58 & 70 & 84 & 102 & 116 & 126 \\ \hline 50 & 50 & 52 & 58 & 69 & 86 & 101 & 120 \\ \hline \end{array} \quad * \quad \begin{array}{|c|c|c|} \hline 0.1 & 0.1 & 0.1 \\ \hline 0.1 & 0.2 & 0.1 \\ \hline 0.1 & 0.1 & 0.1 \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|c|c|c|c|c|} \hline 69 & 95 & 116 & 125 & 129 & 132 \\ \hline 68 & 92 & 110 & 120 & 126 & 132 \\ \hline 66 & 86 & 104 & 114 & 124 & 132 \\ \hline 62 & 78 & 94 & 108 & 120 & 129 \\ \hline 57 & 69 & 83 & 98 & 112 & 124 \\ \hline 53 & 60 & 71 & 85 & 100 & 114 \\ \hline \end{array}$$

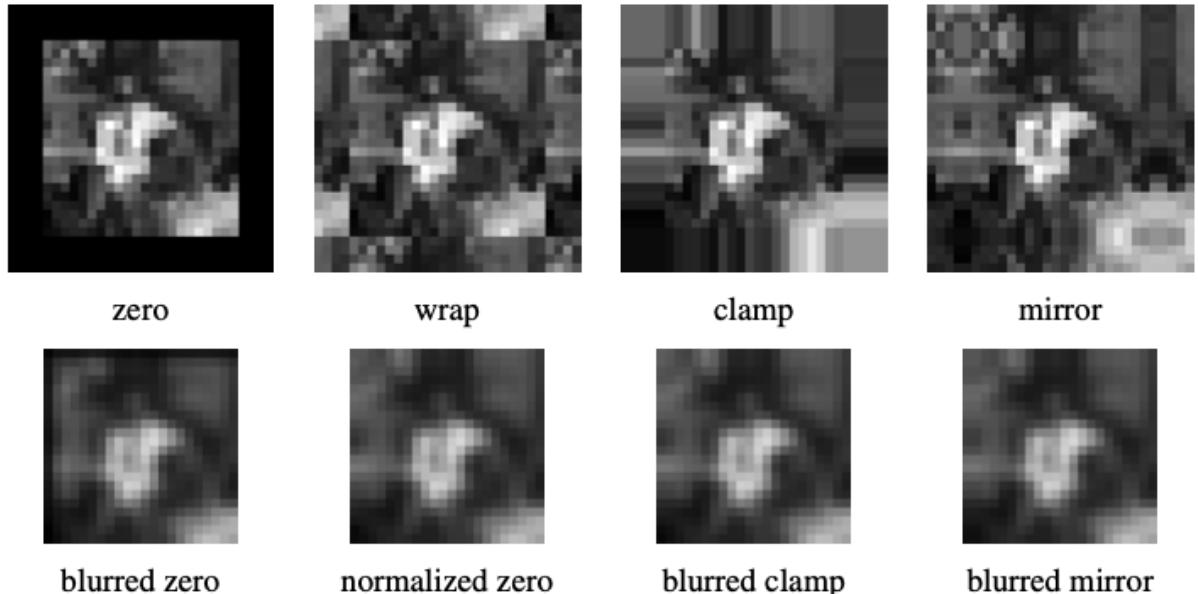
- The correlation/convolution produces a result smaller than the input. Why?
- This may not be desirable in many applications ☹



SOUTH DAKOTA  
STATE UNIVERSITY

# Padding

- To deal with the boundary effects, many padding models have been developed.



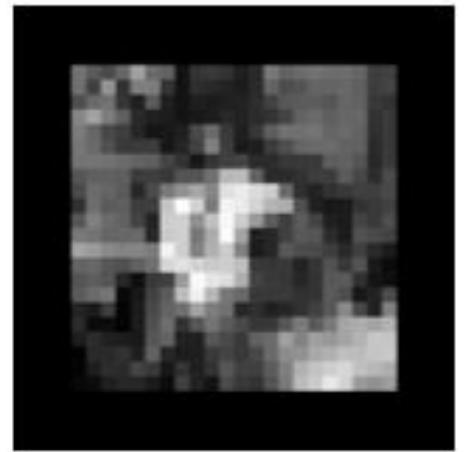
**Figure 3.13** Border padding (top row) and the results of blurring the padded image (bottom row). The normalized zero image is the result of dividing (normalizing) the blurred zero-padded RGBA image by its corresponding soft alpha value.

Reference: Computer Vision: Algorithms and Applications, 2nd ed.

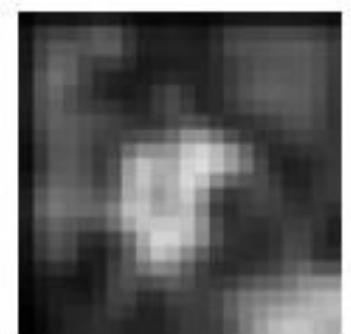
# Zero Padding and Constant Padding

- Zero padding: set all pixels outside the source image to 0
- Usage: simple, common in CNN to keep output size, but can introduce constant-value artifacts.
- Constant padding: set all pixels outside the source image to a specified border value.  
(Zero padding is a special case)

0	0	0	0	0
0	0	0	0	0
0	0	30	60	90
0	0	120	150	180



zero



blurred zero

Reference:

- Computer Vision: Algorithms and Applications, 2nd ed.
- <https://www.mathworks.com/help/images/imfilter-boundary-padding-options.html>

# Clamp (replicate or clamp to edge)

- Definition: repeat edge pixels indefinitely.
- Usage: help reduce boundary artifacts such as edges and contrast compared to padding with zeros or a constant.

30	30	30	60	90
30	30	30	60	90
30	30	30	60	90
120	120	120	150	180



clamp



blurred clamp

Reference:

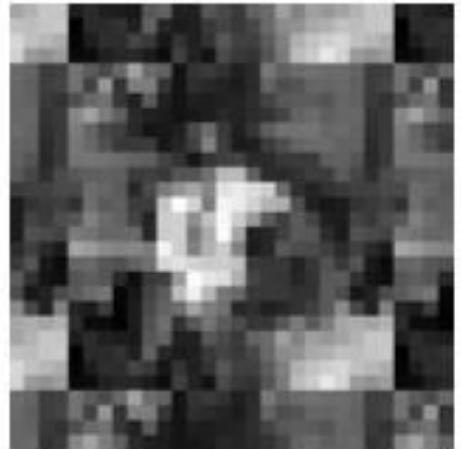
- Computer Vision: Algorithms and Applications, 2nd ed.
- <https://www.mathworks.com/help/images/imfilter-boundary-padding-options.html>

# (Cyclic) Wrap (Repeat or Tile)

- Definition: loop “around” the image in a “toroidal” configuration.
- The padded pixels are copied from the opposite side of the image. The effect gives an appearance of a tiled or looped image.
- Specify this value when you want to treat an image as periodically repeating, such as when you are performing filtering in the frequency domain.

Reference:

- Computer Vision: Algorithms and Applications, 2nd ed.
- <https://www.mathworks.com/help/images/imfilter-boundary-padding-options.html>



wrap

60	90	30	60	90
150	180	120	150	180
60	90	30	60	90
150	180	120	150	180



SOUTH DAKOTA  
STATE UNIVERSITY

# Mirror (or Symmetric)

- Definition: reflect pixels across the image edge.
- Usage: help to remove edge contrast and maintain edge texture.

150	120	120	150	180
60	30	30	60	90
60	30	30	60	90
150	120	120	150	180



mirror



blurred mirror

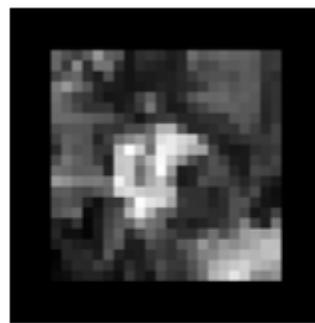
Reference:

- Computer Vision: Algorithms and Applications, 2nd ed.
- <https://www.mathworks.com/help/images/imfilter-boundary-padding-options.html>

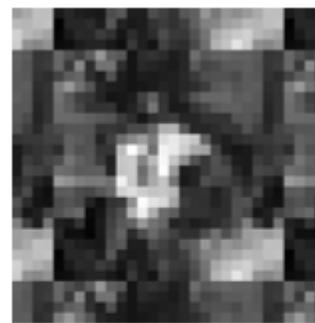


SOUTH DAKOTA  
STATE UNIVERSITY

# Compare blurred results after different paddings



zero



wrap



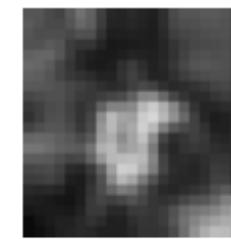
clamp



mirror



blurred zero



normalized zero



blurred clamp



blurred mirror

Observation:

- zero padding darkens the edges,
- clamp (replication) padding propagates border values inward,
- mirror (reflection) padding preserves colors near the borders

**Figure 3.13** Border padding (top row) and the results of blurring the padded image (bottom row). The normalized zero image is the result of dividing (normalizing) the blurred zero-padded RGBA image by its corresponding soft alpha value.

Reference: Computer Vision: Algorithms and Applications, 2nd ed.



SOUTH DAKOTA  
STATE UNIVERSITY

# Animation: Convolution with Zero Padding

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

0	-1	0
-1	5	-1
0	-1	0

114				

Animation source: <https://medium.com/@draj0718/zero-padding-in-convolutional-neural-networks-bf1410438e99>



SOUTH DAKOTA  
STATE UNIVERSITY

Original Image



Erosion



Dilation



Opening



Closing



# Binary Image Processing

Image source: <https://www.geeksforgeeks.org/computer-vision/different-morphological-operations-in-image-processing/>

# Non-Linear Filtering

- The filters we have looked at so far have all been linear, i.e., each output pixel is a weighted summation of neighborhood pixels.
- In many cases, however, better performance can be obtained using a non-linear combination of neighboring pixels.
- Examples: median filtering, bilateral filtering, etc.
- Read *Chapter 3 of Computer Vision: Algorithms and Applications, 2nd ed* if you are interested!



SOUTH DAKOTA  
STATE UNIVERSITY

# Thresholding Operation

- While non-linear filters are often used to enhance grayscale and color images, they are also used extensively to process binary images.
- Such images often occur after a thresholding operation,

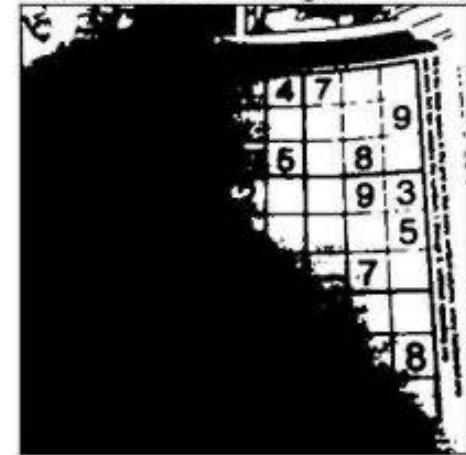
$$\theta(f, t) = \begin{cases} 1 & \text{if } f \geq t, \\ 0 & \text{else,} \end{cases} \quad (3.44)$$

Reference: Computer Vision: Algorithms and Applications, 2nd ed.

Original Image



Global Thresholding (v = 127)



```
# global thresholding  
ret1,th1 = cv.threshold(img,127,255,cv.THRESH_BINARY)
```

Image source:

[https://docs.opencv.org/4.x/d7/d4d/tutorial\\_py\\_thresholding.html](https://docs.opencv.org/4.x/d7/d4d/tutorial_py_thresholding.html)

# Morphological Operations

- Morphological transformations are some simple operations based on the image **shape**.
- It is normally performed on **binary images**.
- It needs two inputs, one is our original image, second one is called **structuring element** or **kernel** which decides the nature of operation.

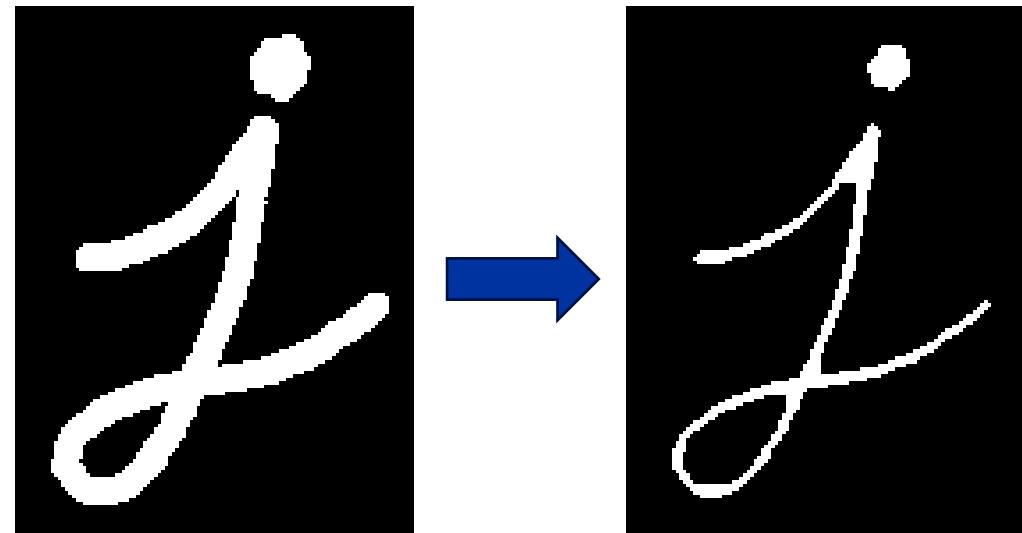
Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)



SOUTH DAKOTA  
STATE UNIVERSITY

# Morphological Operations: Erosion

- The basic idea of erosion is just like soil erosion only, it erodes away the boundaries of foreground object (Always try to keep foreground in white).



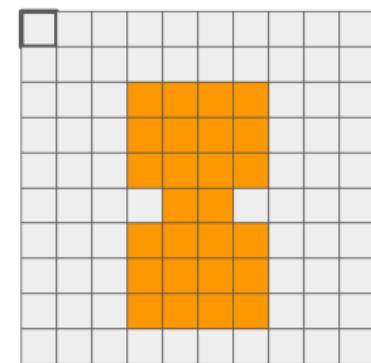
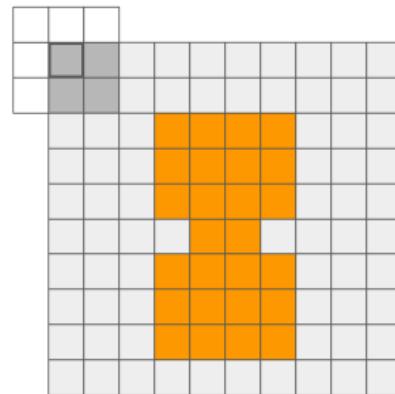
```
import cv2 as cv
import numpy as np

img = cv.imread('j.png', cv.IMREAD_GRAYSCALE)
assert img is not None, "file could not be read, check with os.path.exists()"
kernel = np.ones((5,5),np.uint8)
erosion = cv.erode(img,kernel,iterations = 1)
```

Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

# Morphological Operations: Erosion

- The kernel slides through the image (as in 2D convolution). A pixel in the original image (either 1 or 0) will be considered 1 only if all the pixels under the kernel is 1, otherwise it is eroded (made to zero).

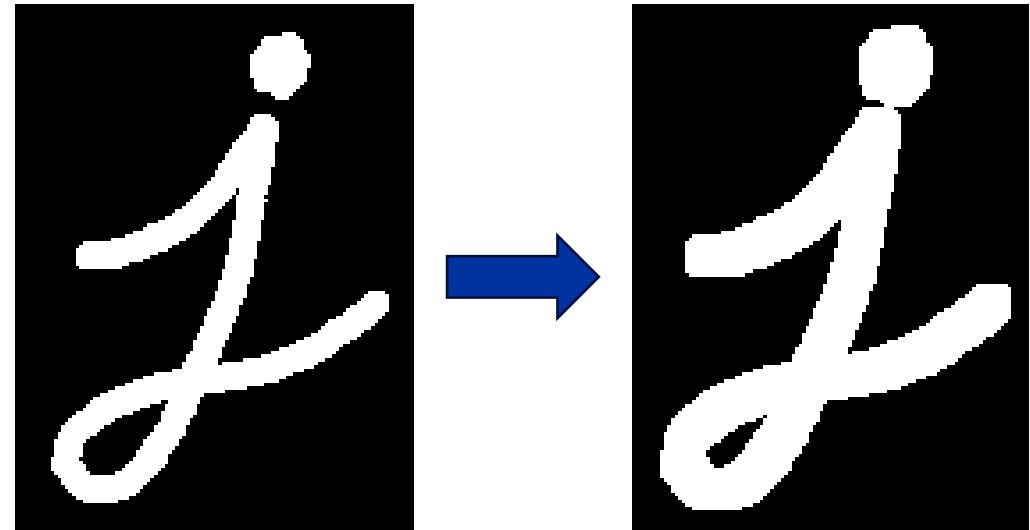


Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

Visualization found: <https://penny-xu.github.io/blog/mathematical-morphology>

# Morphological Operations: Dilation

- It is just opposite of erosion.

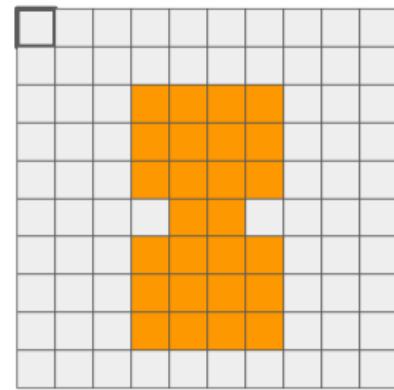
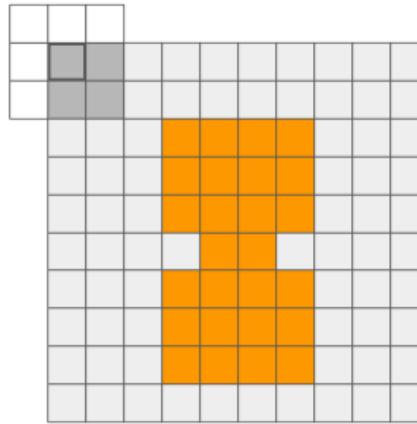


```
dilation = cv.dilate(img,kernel,iterations = 1)
```

Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

# Morphological Operations: Dilation

- Here, a pixel element is '1' if at least one pixel under the kernel is '1'.



Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

Visualization found: <https://penny-xu.github.io/blog/mathematical-morphology>

# Morphological Operations: Opening

- Opening is just another name of **erosion followed by dilation**.
- It is useful in removing noise.



```
opening = cv.morphologyEx(img, cv.MORPH_OPEN, kernel)
```

Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

# Morphological Operations: Closing

- Closing is reverse of **Opening, Dilation followed by Erosion.**
- It is useful in closing small holes inside the foreground objects, or small black points on the object.



```
closing = cv.morphologyEx(img, cv.MORPH_CLOSE, kernel)
```

Reference: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

# Morphology for Data Preprocessing

Manually labelled segmentation masks may contain noises. We can consider using morphological operations to preprocessing the annotated masks before training segmentation models.



Examples from  
Linemod Occlusion  
Dataset.

<https://bop.felk.cvut.cz/datasets/#LM-O>



SOUTH DAKOTA  
STATE UNIVERSITY



SOUTH DAKOTA  
STATE UNIVERSITY



# Why Features Matter in Computer Vision?

Image source: [https://docs.opencv.org/4.x/da/df5/tutorial\\_py\\_sift\\_intro.html](https://docs.opencv.org/4.x/da/df5/tutorial_py_sift_intro.html)

# Let's Play a Game: Jigsaw Puzzle!

- Link:  
<https://onlinejigsawpuzzles.net/puzzle.php?image=images/puzzle/Spring-in-farmland-Sweden-old-bar--in-rural-surroundings.jpg#>
- Note: you can change the number of pieces
- Think: How do you solve the puzzle? What visual cues help you place pieces?



SOUTH DAKOTA  
STATE UNIVERSITY

# Jigsaw Puzzles in Computer Vision Research

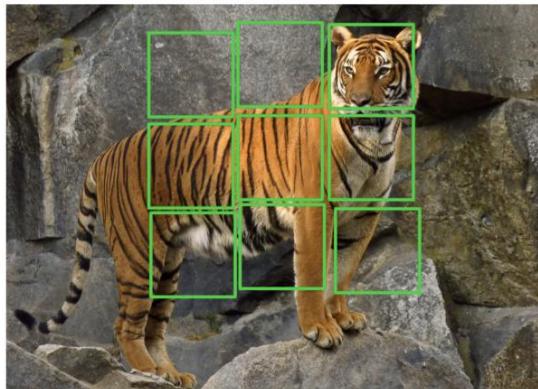
- In a jigsaw puzzle, we reconstruct a global image from local pieces.
- During the process, we may consider edge continuity, texture similarity, color consistency, etc.
- Computer vision aims to automate the same reasoning process.
- Example research papers: Unsupervised Learning of Visual Representations by Solving Jigsaw Puzzles (ECCV 2016, paper link: <https://arxiv.org/abs/1603.09246> )



Screenshot for

<https://onlinejigsawpuzzles.net/puzzle.php?image=images/puzzle/Spring-in-farmland-Sweden-old-bar--in-rural-surroundings.jpg#>

# Jigsaw Puzzles in Computer Vision Research



(a)



(b)



(c)

Fig. 1: Learning image representations by solving Jigsaw puzzles. (a) The image from which the tiles (marked with green lines) are extracted. (b) A puzzle obtained by shuffling the tiles. Some tiles might be directly identifiable as object parts, but others are ambiguous (*e.g.*, have similar patterns) and their identification is much more reliable when all tiles are jointly evaluated. In contrast, with reference to (c), determining the relative position between the central tile and the top two tiles from the left can be very challenging [10].

Example research papers:  
Unsupervised Learning of Visual Representations by Solving Jigsaw Puzzles (ECCV 2016, paper link: <https://arxiv.org/abs/1603.09246>)

# What is a Feature in Computer Vision?

- A feature in computer vision is an identifiable part of an image or video that conveys meaningful information for tasks like object detection, tracking, or classification.
- Features are often distinct patterns, such as edges, corners, textures, or specific shapes, that algorithms use to understand and analyze visual data.

Texts adapted from <https://milvus.io/ai-quick-reference/what-is-a-feature-in-computer-vision>



Screenshot for  
<https://onlinejigsawpuzzles.net/puzzle.php?image=images/puzzle/Spring-in-farmland-Sweden-old-bar-in-rural-surroundings.jpg#>

# Why Not Just Use Pixels?

- Pixels are sensitive to noise, lighting variations, viewpoint changes, etc.
- Therefore, pixels are not stable representations for robust computer vision tasks.

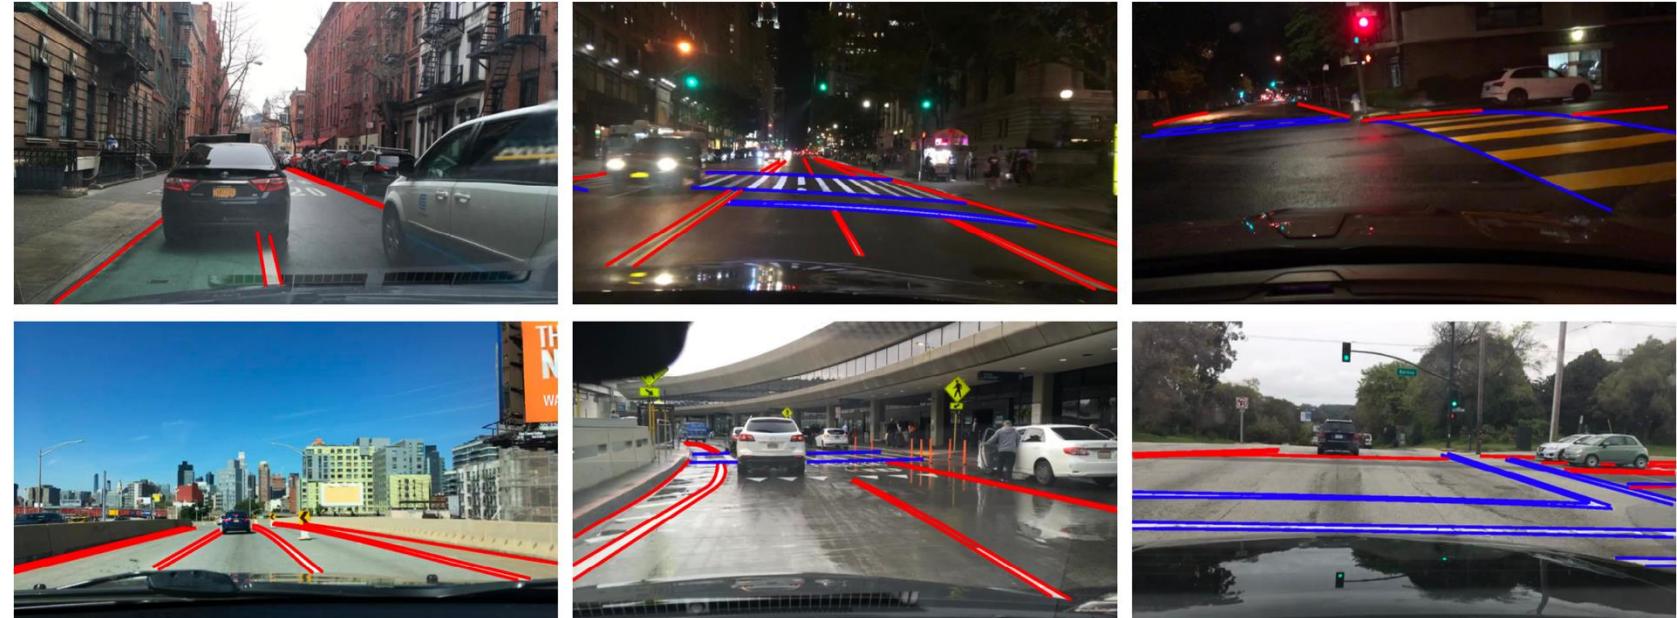


Image source: BDD100K: A Large-scale Diverse Driving Video Database  
<https://bair.berkeley.edu/blog/2018/05/30/bdd/>

# Classical (Hand-Crafted) Features

- Traditional feature extraction methods rely on mathematical techniques to identify and describe these key points.
- Algorithms like SIFT (Scale-Invariant Feature Transform) detect stable features by analyzing gradients, corners, or blobs in an image.

Texts adapted from <https://milvus.io/ai-quick-reference/what-is-a-feature-in-computer-vision>



Image found at  
<https://ics.uci.edu/~majumder/VC/211HW3/vlfeat/doc/overview/sift.html>

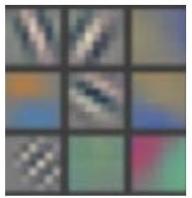
# Learned Features

- Modern approaches, particularly deep learning, automate feature extraction using convolutional neural networks (CNNs).
- In CNNs, layers learn hierarchical features directly from data.
- Early layers detect simple patterns like edges or color gradients, while deeper layers combine these to recognize complex shapes or objects.

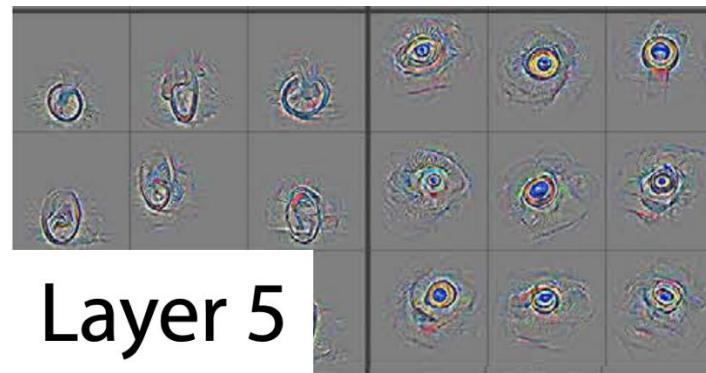
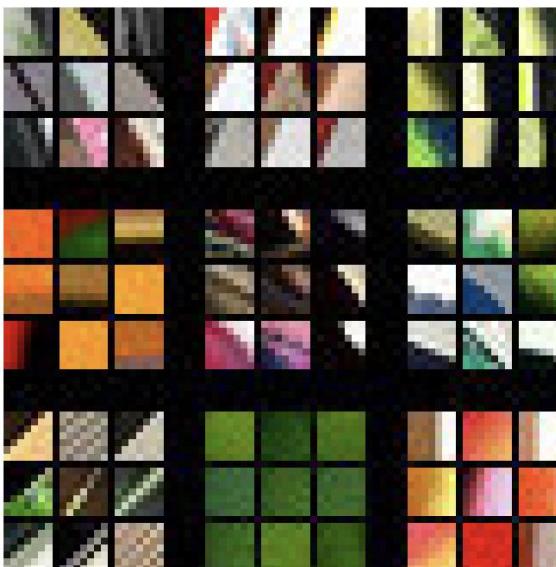


SOUTH DAKOTA  
STATE UNIVERSITY

# “Visualizing and Understanding Convolutional Networks” (ECCV 2014)



Layer 1



Layer 5



Check full Figure 2 in the original paper.  
You can read the paper in this link:  
<https://arxiv.org/pdf/1311.2901>



SOUTH DAKOTA  
STATE UNIVERSITY



SOUTH DAKOTA  
STATE UNIVERSITY

# Feature Detection I: Edge Detection

Credits: some mathematical formulations taken from [http://vision.stanford.edu/teaching/cs131\\_fall1718/files/05\\_edges.pdf](http://vision.stanford.edu/teaching/cs131_fall1718/files/05_edges.pdf) with some modifications! Because this document explicitly differentiate convolution and correlation.

Image source: [https://docs.opencv.org/4.x/d4/d22/tutorial\\_py\\_canny.html](https://docs.opencv.org/4.x/d4/d22/tutorial_py_canny.html)

Original Image



Edge Image



# Edges are about changes...

- Derivatives in 1D

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x$$

- Discrete derivative in 1D

$$\frac{df}{dx} = \frac{f(x) - f(x - 1)}{1} = f'(x)$$

$$\frac{df}{dx} = f(x) - f(x - 1) = f'(x)$$



SOUTH DAKOTA  
STATE UNIVERSITY

# Types of Discrete derivative in 1D (Approximation!)

**Backward**

$$\frac{df}{dx} = f(x) - f(x-1) = f'(x)$$

**Forward**

$$\frac{df}{dx} = f(x) - f(x+1) = f'(x)$$

**Central**

$$\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$$



SOUTH DAKOTA  
STATE UNIVERSITY

# 1D discrete derivate filters

Input:  $[f(x-1), f(x), f(x+1)]$

- Backward filter:

$$f(x) - f(x-1) = f'(x)$$

Kernel: [-1 1 0]

- Forward:

$$f(x) - f(x+1) = f'(x)$$

Kernel: [0 1 -1]

- Central:

$$f(x+1) - f(x-1) = f'(x)$$

Kernel: [-1 0 1]

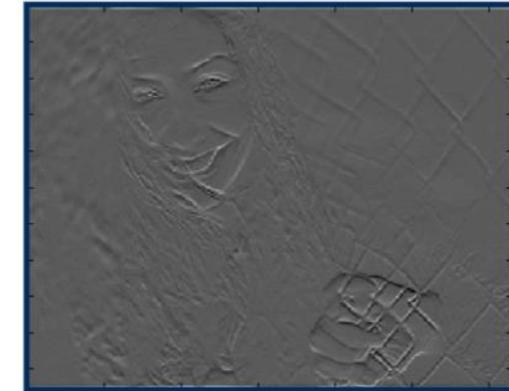


SOUTH DAKOTA  
STATE UNIVERSITY

# 3x3 image gradient filters

$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



SOUTH DAKOTA  
STATE UNIVERSITY

Visualization found in  
[http://vision.stanford.edu/teaching/cs131\\_fall1718/files/05\\_edges.pdf](http://vision.stanford.edu/teaching/cs131_fall1718/files/05_edges.pdf)

# Edge Detection – Finite Difference Filter

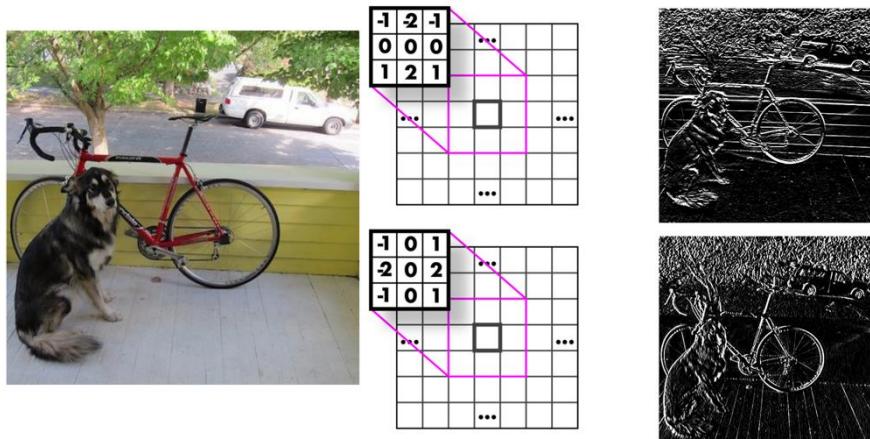
- **Convolution Mask:**

Prewitt

$$\frac{\partial I}{\partial x} \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$
$$\frac{\partial I}{\partial y} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$$

Sobel

$$\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$
$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$



Note: when a convolution kernel sums to 1 (e.g., box filter), it guarantees the overall brightness of the image stays the same. This is desired for many image processing operations.  
For Prewitt and Sobel kernels, the sum is 0 → detect changes



SOUTH DAKOTA  
STATE UNIVERSITY

# Discrete derivate in 2D

Given function

$$f(x, y)$$

Gradient vector

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} \\ \frac{\partial f(x, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

Gradient magnitude

$$|\nabla f(x, y)| = \sqrt{f_x^2 + f_y^2}$$

Gradient direction

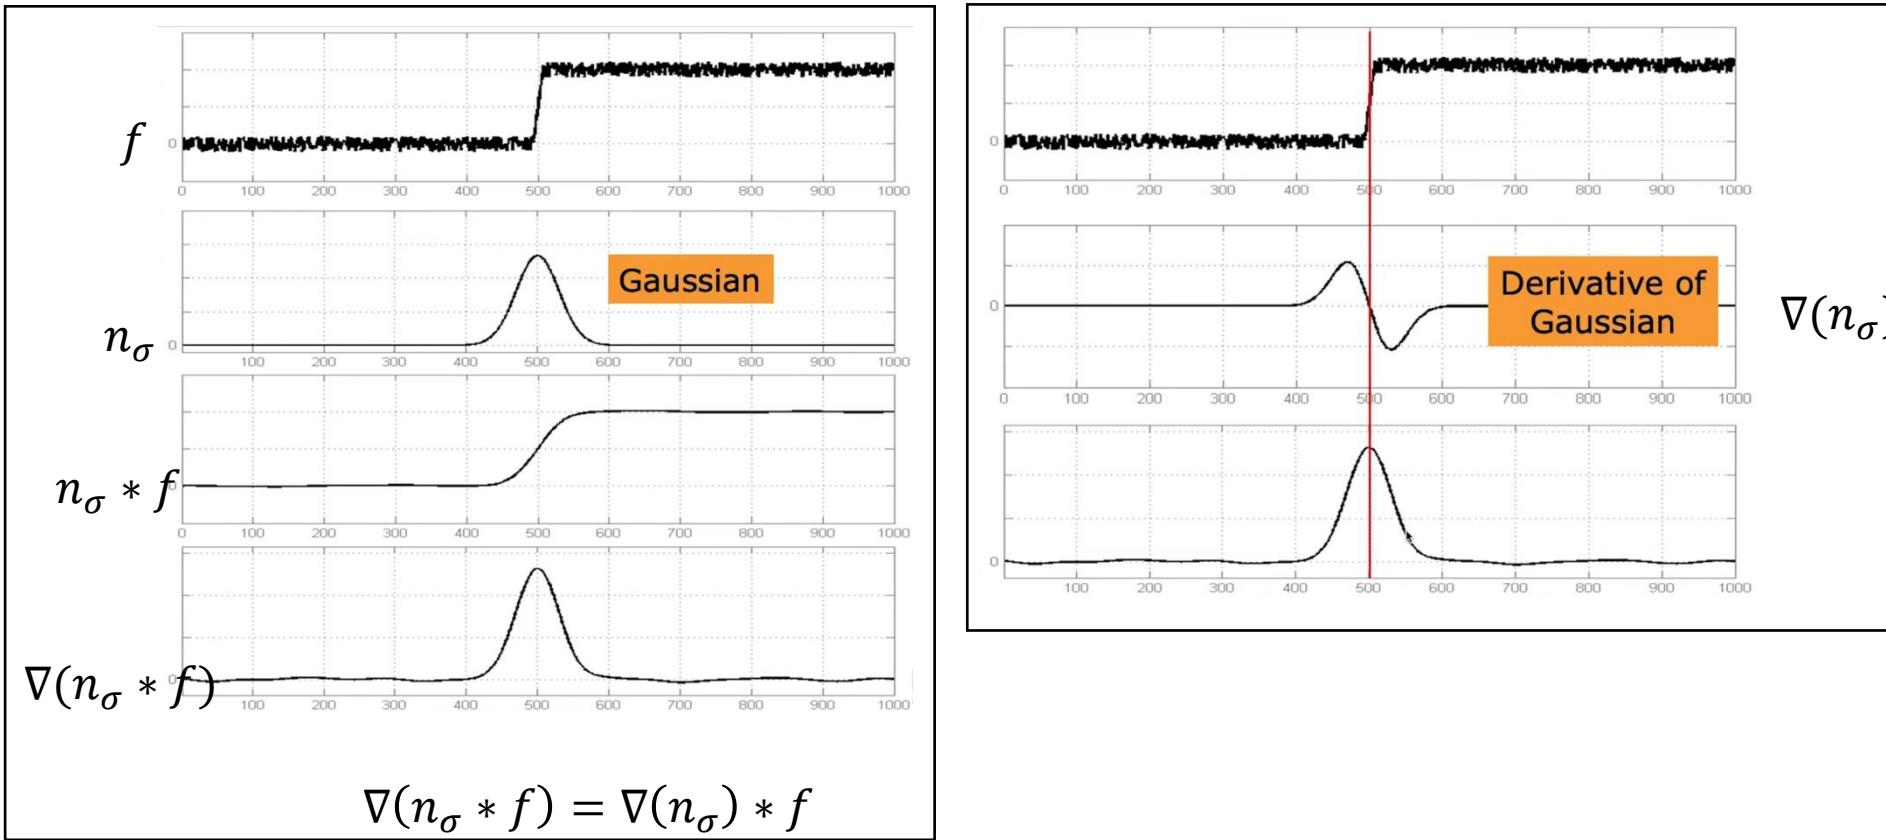
$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$



SOUTH DAKOTA  
STATE UNIVERSITY

# Derivative of Gaussian Filter (DoG)

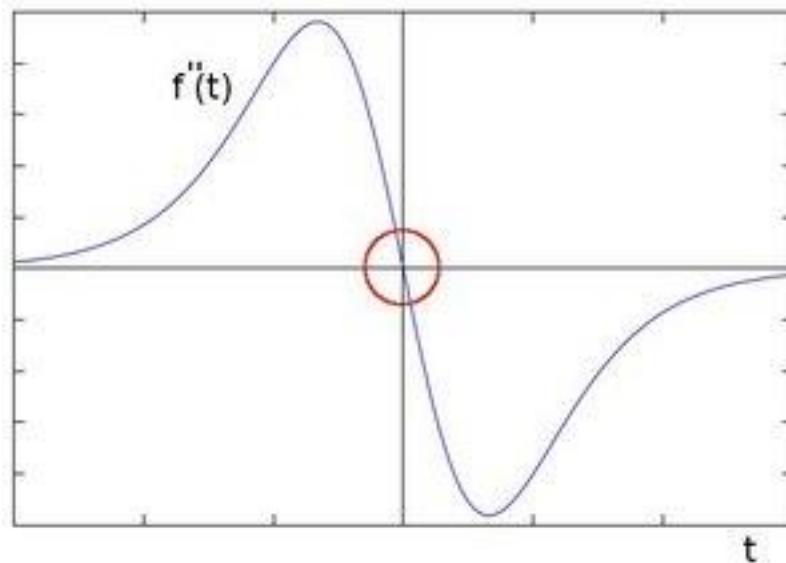
- *Idea.* Smooth before differentiation!



SOUTH DAKOTA  
STATE UNIVERSITY

# Laplacian Filter (Based on 2<sup>nd</sup> derivative)

1D intuition: zero crossing



2D Laplacian Operator

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

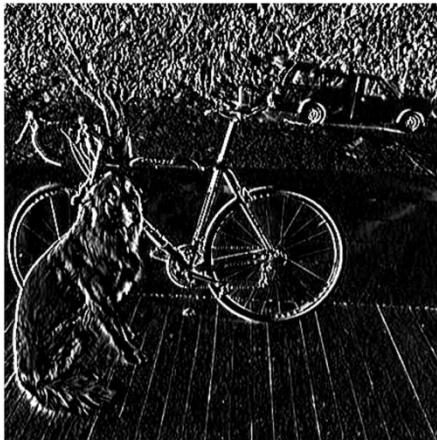


0	1	0
1	-4	1
0	1	0

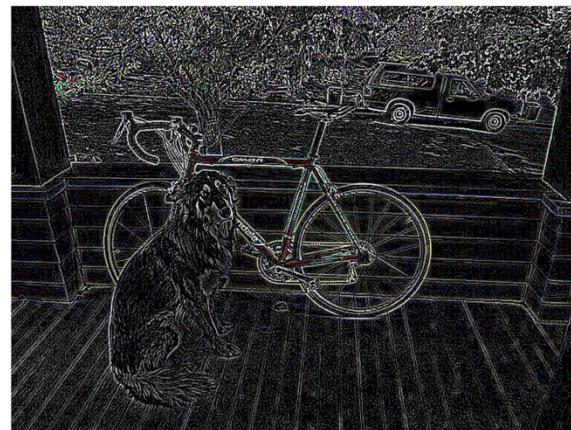
Reference: [https://docs.opencv.org/3.4/d5/db5/tutorial\\_laplace\\_operator.html](https://docs.opencv.org/3.4/d5/db5/tutorial_laplace_operator.html)

# Summary: Edge Detection

Prewitt			Sobel		
$\frac{\partial I}{\partial x}$	-1 0 1 -1 0 1 -1 0 1		-1 0 1 -2 0 2 -1 0 1		
$\frac{\partial I}{\partial y}$	1 1 1 0 0 0 -1 -1 -1		1 2 1 0 0 0 -1 -2 -1		



Laplacian		
1	4	1
4	-20	4
1	1	1



Gradient	Laplacian
Provides <b>location</b> , <b>magnitude</b> , and <b>direction</b> of the edge	Provides only <b>location</b> of the edge
Detection using Maxima Thresholding	Detection based on Zero-crossing
<ul style="list-style-type: none"><li>- <b>Non-linear</b> operation</li><li>- Requires two convolutions</li></ul>	<ul style="list-style-type: none"><li>- <b>Linear</b> operation</li><li>- Requires only one convolution</li></ul>



SOUTH DAKOTA  
STATE UNIVERSITY

# Canny Edge Detection

Credit: slides of this section are adapted from CSC 492/592, Spring 2024

Image source: [https://docs.opencv.org/4.x/da/d22/tutorial\\_py\\_canny.html](https://docs.opencv.org/4.x/da/d22/tutorial_py_canny.html)

Original Image

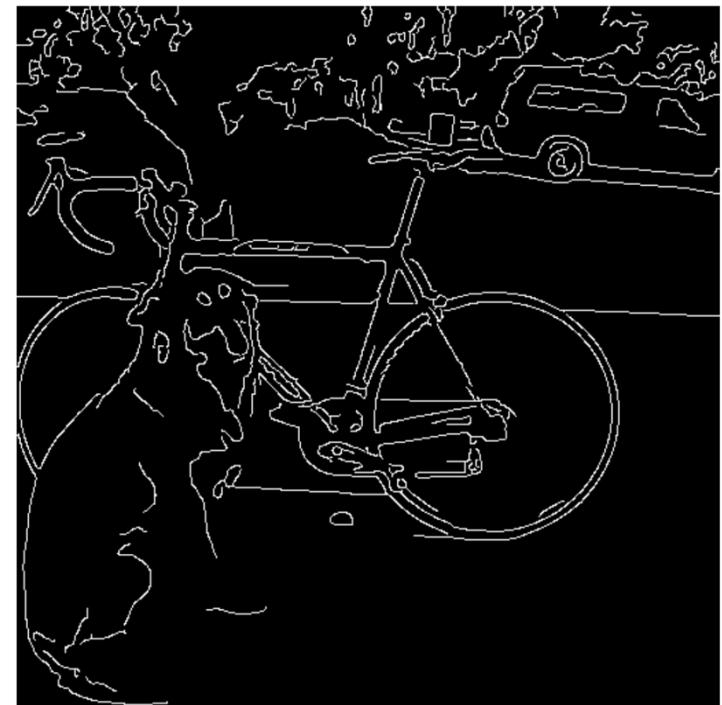


Edge Image



# Canny Edge Detection

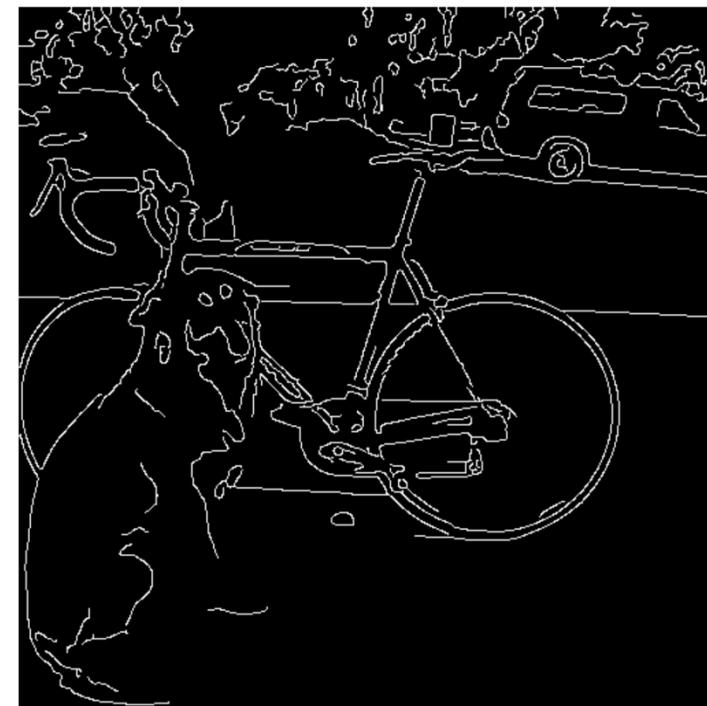
- The Canny edge detector is an edge detection operator that uses a multi-stage algorithm to detect a wide range of edges in images.
- It was developed by John F. Canny in 1986.
- Among the edge detection methods developed so far, Canny's algorithm is one of the most strictly defined methods that provides good and reliable detection.



Reference: [https://en.wikipedia.org/wiki/Canny\\_edge\\_detector](https://en.wikipedia.org/wiki/Canny_edge_detector)

# Canny Edge Detection: Algorithm

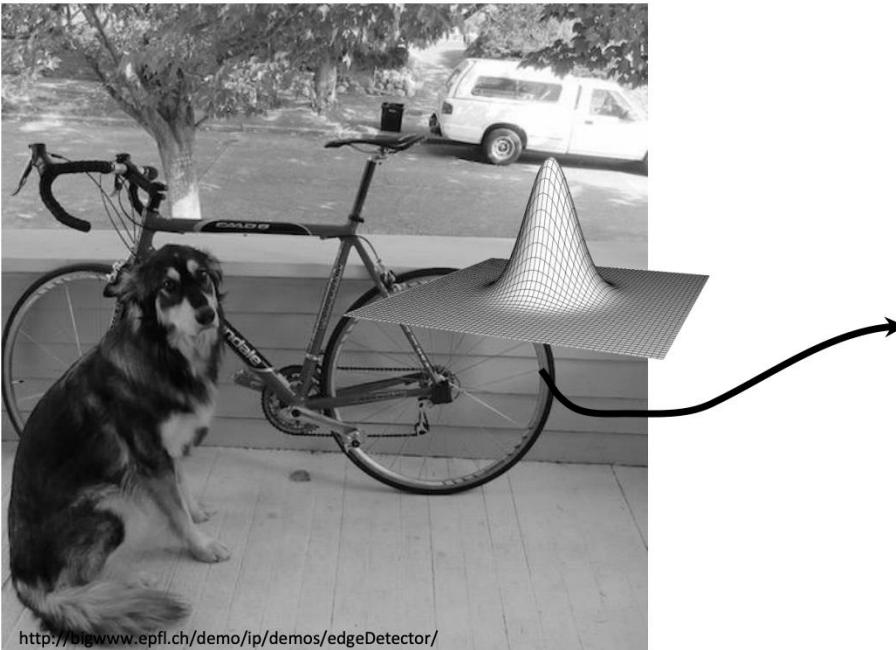
1. Smooth image
2. Calculate gradient direction and magnitude
3. Perform non-maximum suppression
4. Threshold the edges
5. Link the edges



SOUTH DAKOTA  
STATE UNIVERSITY

# Step 1: Smooth image

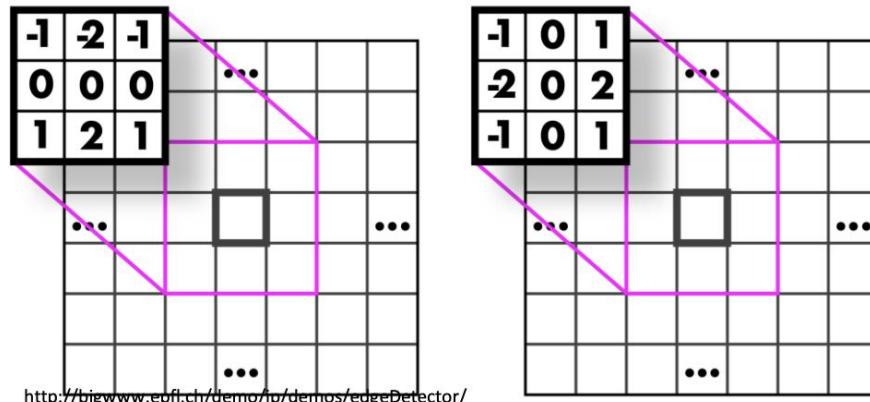
Apply Gaussian filter to smooth the image to remove the noise.



SOUTH DAKOTA  
STATE UNIVERSITY

# Step 2: Calculate gradient direction and magnitude

	Prewitt	Sobel																		
$\frac{\partial I}{\partial x}$	<table border="1"><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1	-1	0	1	-1	0	1	<table border="1"><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-2</td><td>0</td><td>2</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1	-2	0	2	-1	0	1
-1	0	1																		
-1	0	1																		
-1	0	1																		
-1	0	1																		
-2	0	2																		
-1	0	1																		
$\frac{\partial I}{\partial y}$	<table border="1"><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-1</td><td>-1</td></tr></table>	1	1	1	0	0	0	-1	-1	-1	<table border="1"><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-2</td><td>-1</td></tr></table>	1	2	1	0	0	0	-1	-2	-1
1	1	1																		
0	0	0																		
-1	-1	-1																		
1	2	1																		
0	0	0																		
-1	-2	-1																		



SOUTH DAKOTA  
STATE UNIVERSITY

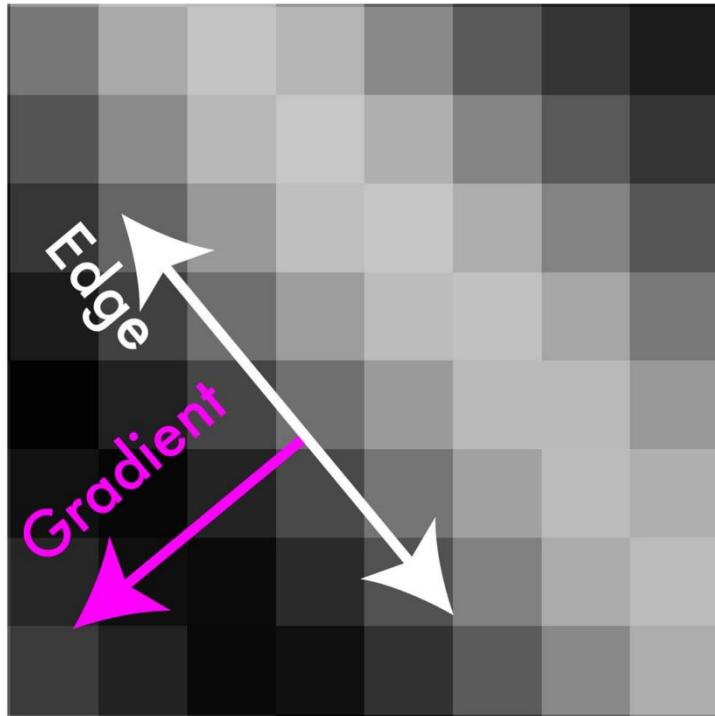
## Step 3: Perform non-maximum suppression

- **Why?** Want single pixel edges, not thick blurry lines
- **Idea.** Check nearby pixels and keep the high response
- **How?** Go through all the points on the gradient intensity matrix and finds the pixels with the maximum value in the **edge directions**



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression

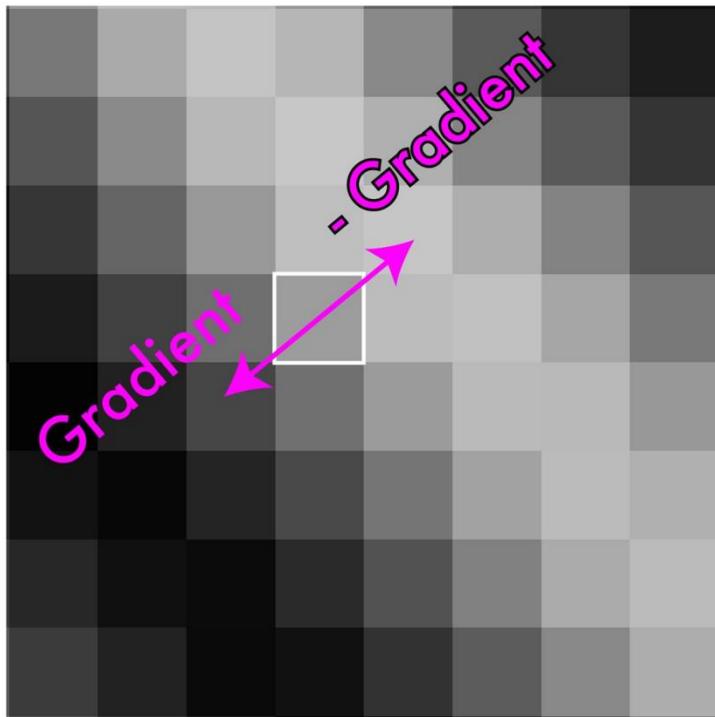


Go through all the points on the gradient intensity matrix and finds the pixels with the maximum value in the edge directions.



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression



Example pixel 1:

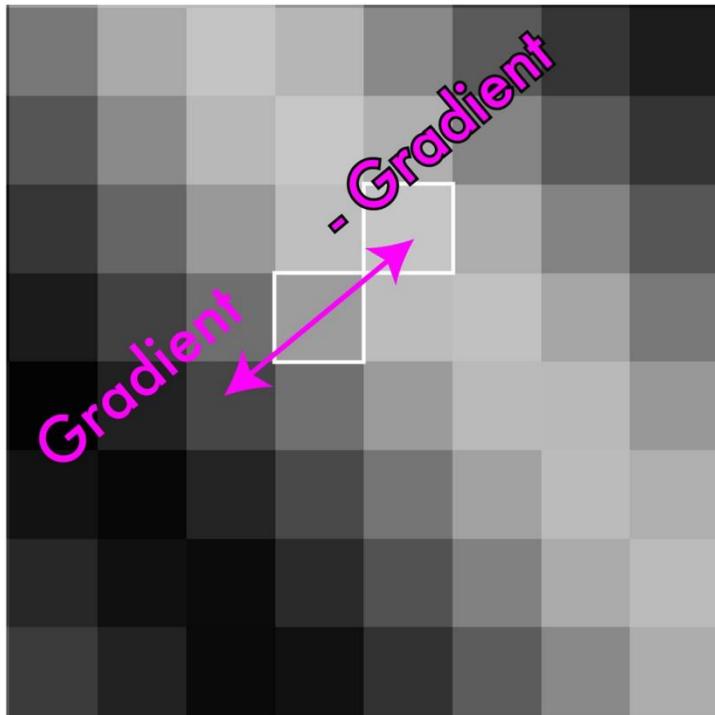
(highlighted in white box)

compare its gradient magnitude with  
the magnitudes of its neighboring  
pixels along the gradient direction



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression



Example pixel 1:

(highlighted in left-bottom white box)

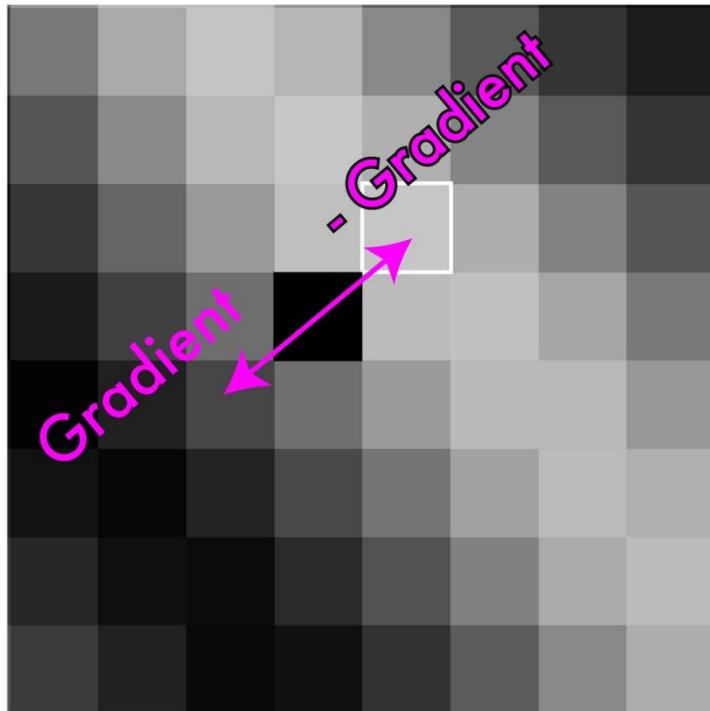
compare its gradient magnitude with  
the magnitudes of its neighboring  
pixels along the gradient direction

→ We find a neighboring pixel with  
larger gradient magnitude



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression



Example pixel 1:

(highlighted in left-bottom white box)

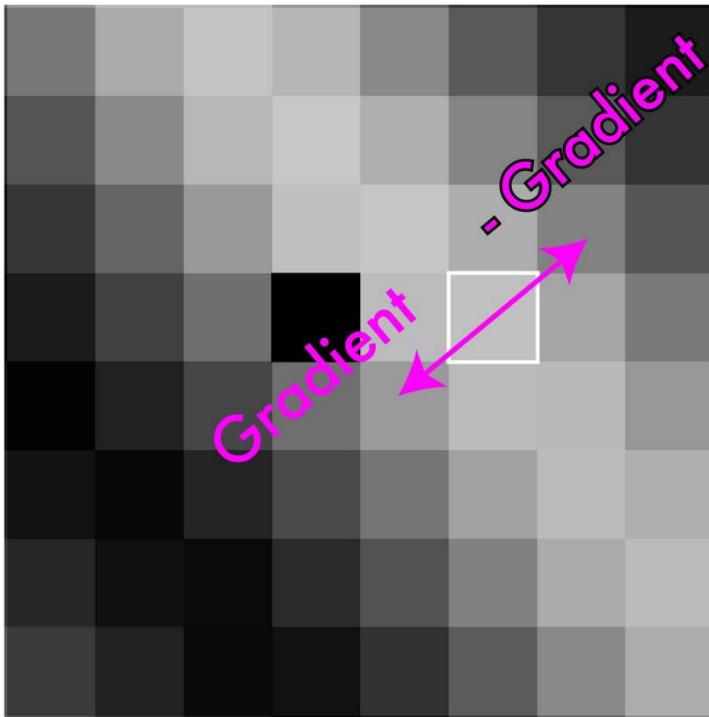
compare its gradient magnitude with  
the magnitudes of its neighboring  
pixels along the gradient direction

- We find a neighboring pixel with  
larger gradient magnitude
- It is suppressed ☹



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression



Example pixel 2:

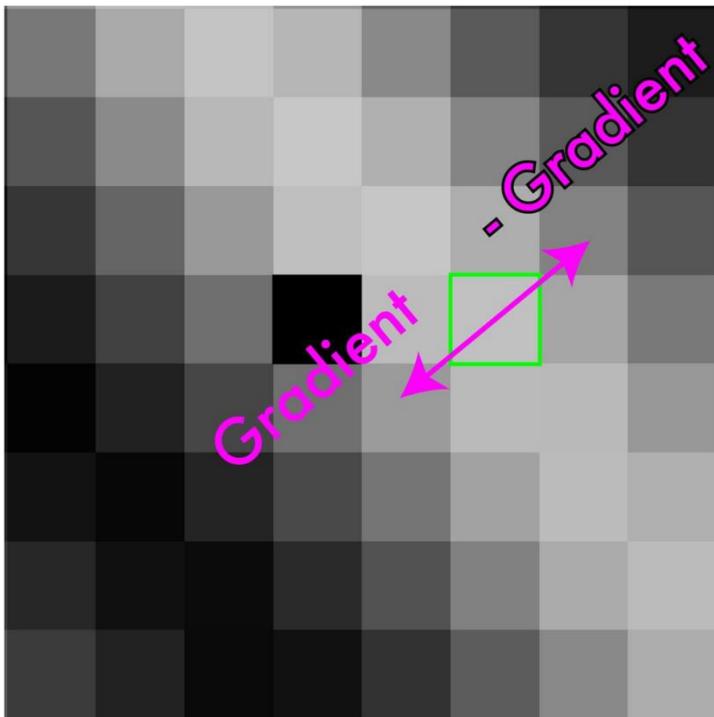
(highlighted in white box)

compare its gradient magnitude with  
the magnitudes of its neighboring  
pixels along the gradient direction



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression



Example pixel 2:

(highlighted in green box now)

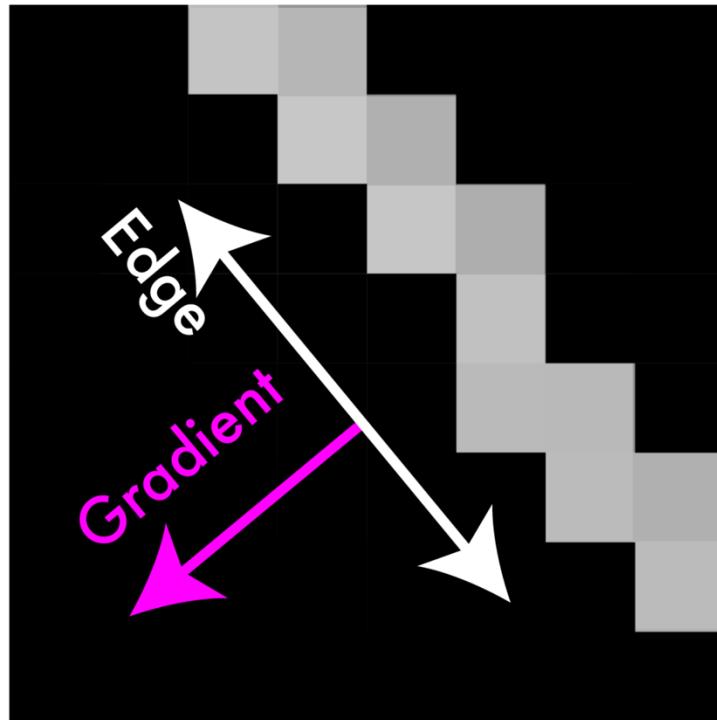
compare its gradient magnitude with the magnitudes of its neighboring pixels along the gradient direction

- Greater than that of both of its neighboring pixels
- It is kept ☺



SOUTH DAKOTA  
STATE UNIVERSITY

## Step 3: Perform non-maximum suppression

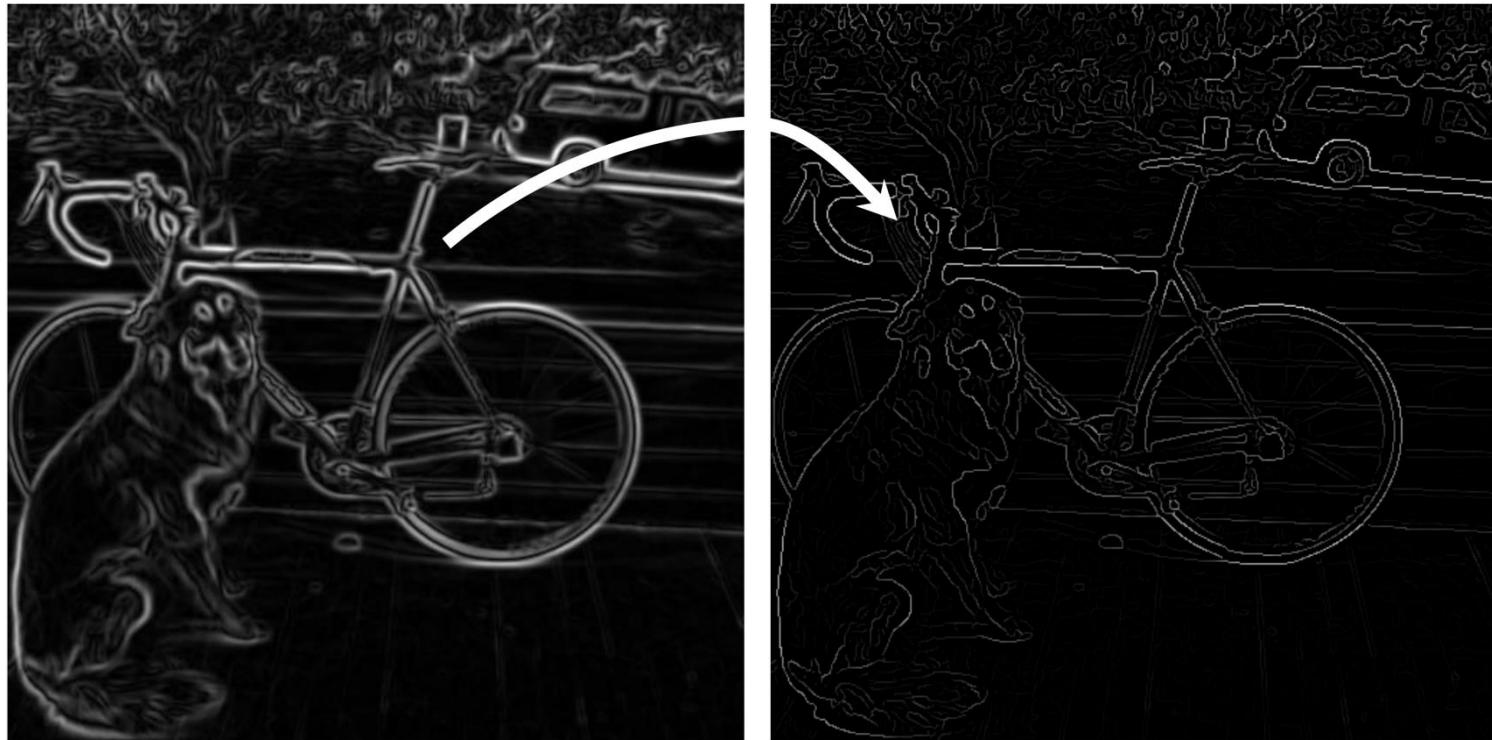


This process is performed for all pixels in the image, resulting in a thinned set of edge pixels where only local maxima **along the edges** remain.



SOUTH DAKOTA  
STATE UNIVERSITY

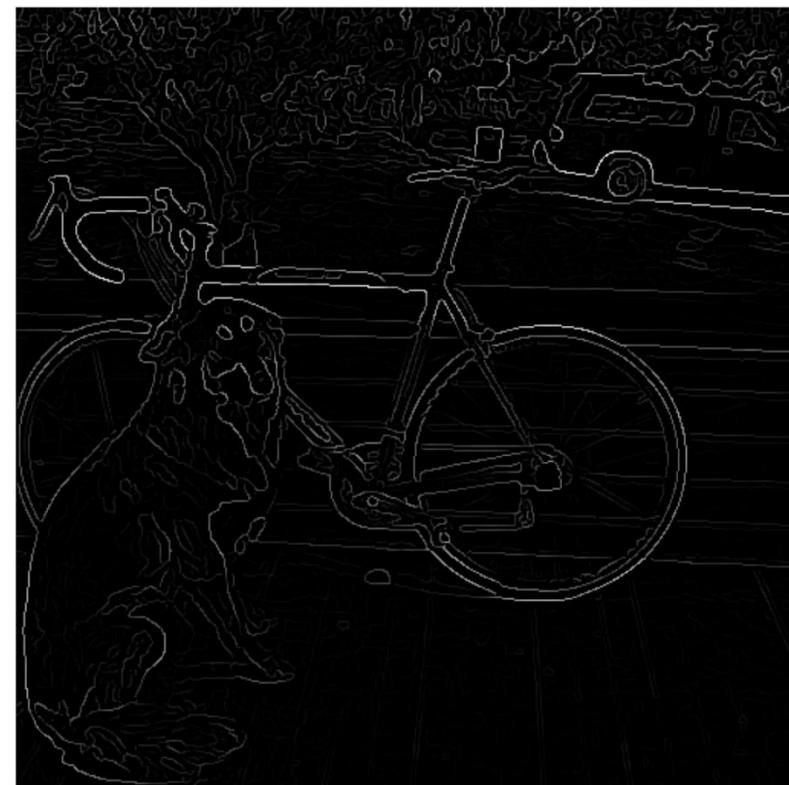
## Step 3: Perform non-maximum suppression



SOUTH DAKOTA  
STATE UNIVERSITY

# Step 4: Threshold the edges

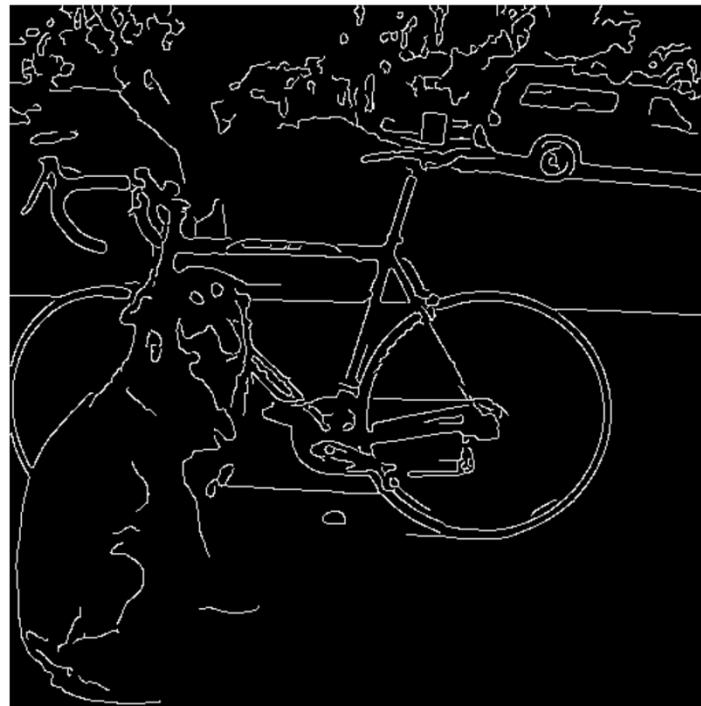
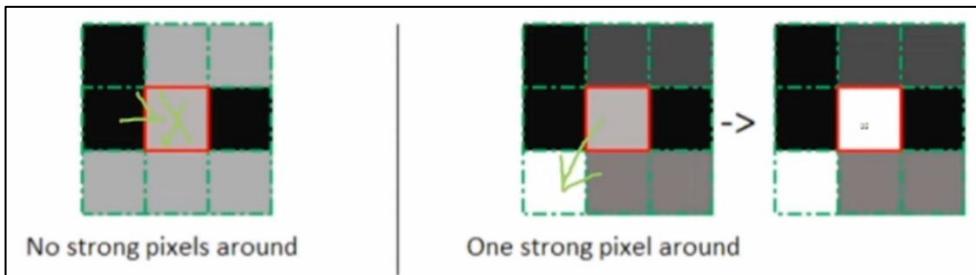
- **Why?** Still some noise
- **Idea.** Thresholding edges
- **How?**
  - Categorize edges into **strong**, **weak**, **non-relevant**
  - 2 thresholds  $T$  and  $t$ , 3 cases:
    - .  $T < \cdot$  : strong edge
    - .  $t < \cdot < T$  : weak edge
    - .  $\cdot < t$  : non-relevant



SOUTH DAKOTA  
STATE UNIVERSITY

# Step 5: Link the edges

- **Idea.**
  - Strong edges are edges!
  - Weak edges are edges if and only if they are connected to strong edges
- **How?**
  - Look in some neighborhood (usually 8 closet)



SOUTH DAKOTA  
STATE UNIVERSITY

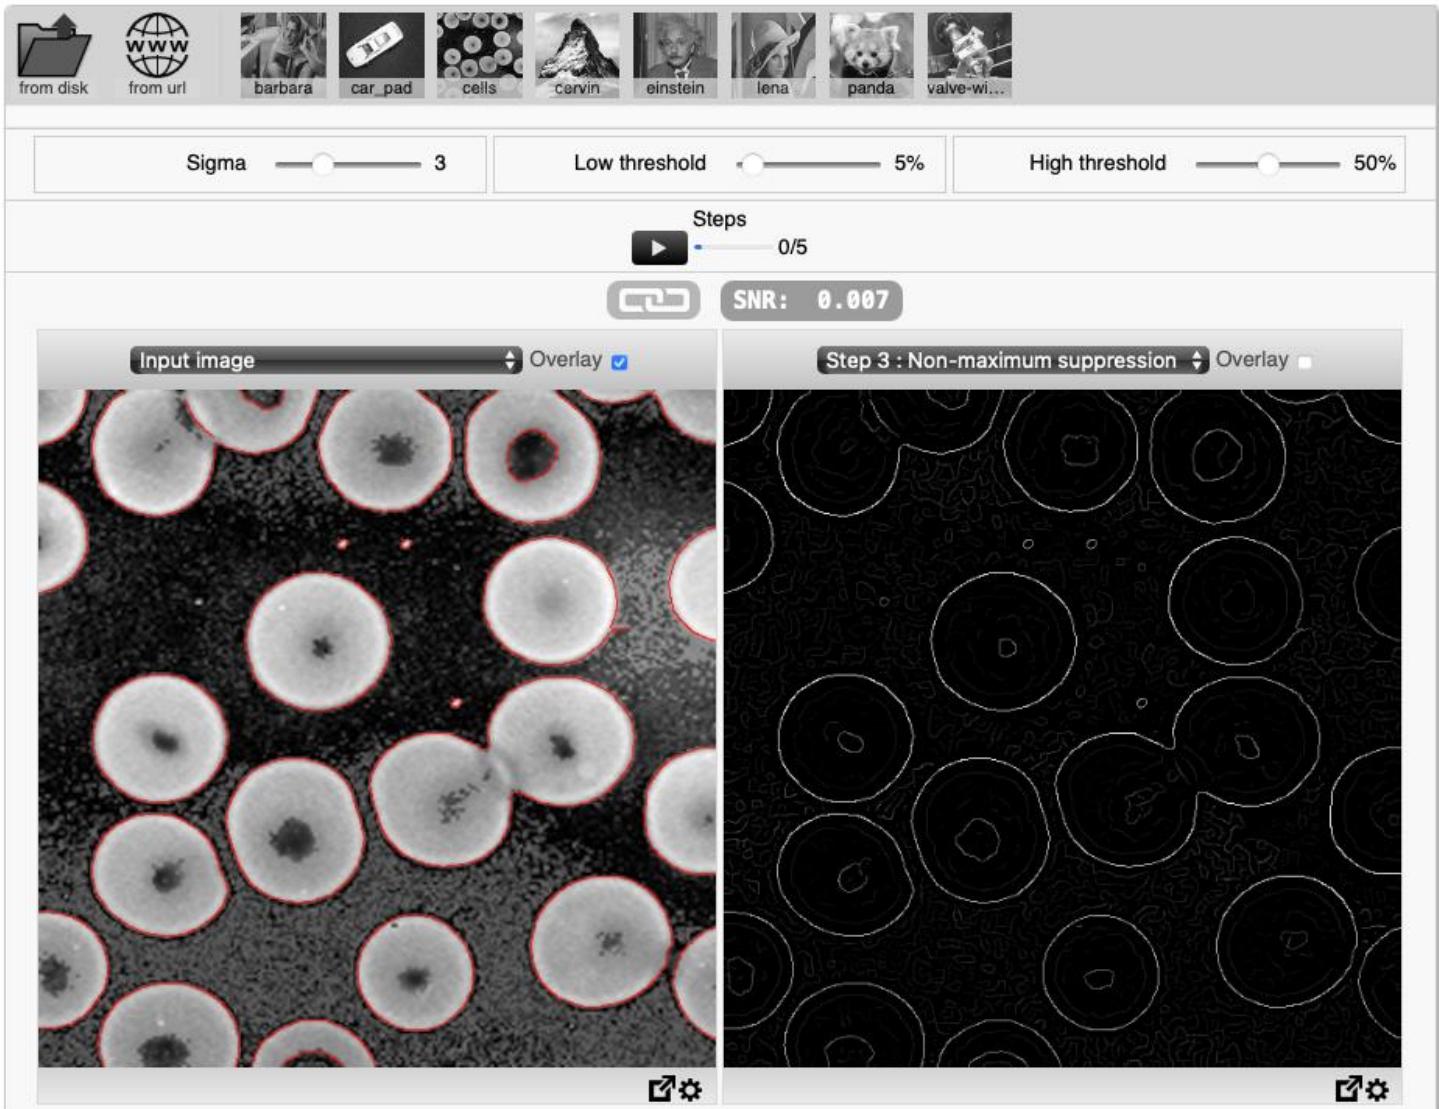
# Demo

Let's try:

<https://bigwww.epfl.ch/demo/ip/demos/edgeDetector/>

Think: how the hyperparameters affect the results?

In OpenCV: simply call “cv2.Canny” function  
(reference:  
[https://docs.opencv.org/4.x/d22/tutorial\\_py\\_canny.html](https://docs.opencv.org/4.x/d22/tutorial_py_canny.html) )



SOUTH DAKOTA  
STATE UNIVERSITY

# Takeaway

- Padding?
- Binary Image: Erosion, Dilation?
- Why Features matter in computer vision?
- Different ways to detect edges
  - Prewitt & Sobel
  - DoG
  - Laplacian
  - Canny Edge Detector



Questions?



SOUTH DAKOTA  
STATE UNIVERSITY